

A Relational Separation Logic for Effect Handlers

PAULO EMÍLIO DE VILHENA, Imperial College London, United Kingdom

SIMCHA VAN COLLEM, Radboud University Nijmegen, The Netherlands

INES WRIGHT, Aarhus University, Denmark

ROBBERT KREBBERS, Radboud University Nijmegen, The Netherlands

Effect handlers offer a powerful and relatively simple mechanism for controlling a program's flow of execution. Since their introduction, an impressive array of verification tools for effect handlers has been developed. However, to this day, no framework can express and prove *relational properties* about programs that use effect handlers in languages such as OCaml and Links, where programming features like mutable state and concurrency are readily available. To this end, we introduce *blaze*, the first *relational separation logic* for effect handlers. We build *blaze* on top of the Iris framework for concurrent separation logic in Rocq, thereby enjoying the rigour of a mechanised theory and all the reasoning properties of a modern fully-fledged concurrent separation logic, such as modular reasoning about stateful concurrent programs and the ability to introduce user-defined ghost state. In addition to familiar reasoning rules, such as the bind rule and the frame rule, *blaze* offers rules to reason modularly about programs that perform and handle effects. Significantly, when verifying that two programs are related, *blaze* *does not* require that effects and handlers from one program be in correspondence with effects and handlers from the other. To assess this flexibility, we conduct a number of case studies: most noticeably, we show how different implementations of an asynchronous-programming library using effects are related to *truly concurrent* implementations. As side contributions, we introduce two new, simple, and general reasoning rules for concurrent relational separation logic that are independent of effects: a *logical-fork rule* that allows one to reason about an arbitrary program phrase as if it had been spawned as a thread and a *thread-swap rule* that allows one to reason about how threads are scheduled.

CCS Concepts: • Theory of computation → Control primitives; Concurrency; Separation logic.

Additional Key Words and Phrases: Effect Handlers, Concurrency, Relational Separation Logic, Iris, Rocq

ACM Reference Format:

Paulo Emílio de Vilhena, Simcha van Collem, Ines Wright, and Robbert Krebbers. 2026. A Relational Separation Logic for Effect Handlers. *Proc. ACM Program. Lang.* 10, POPL, Article 34 (January 2026), 29 pages. <https://doi.org/10.1145/3776676>

1 Introduction

Effect handlers [45] are a powerful programming abstraction that separates the use of an effect from its implementation, allowing programmers to write effectful code independently of how these effects are implemented. Its programming interface offers the ability to *perform* and to *handle* effects. Performing an effect is similar to raising an exception: execution is suspended and control is transferred to an enclosing pre-installed handler. Handling an effect is also similar to handling an exception with the key difference that, in addition to the effect's payload, the effect handler also has access to a first-class representation of the suspended program, a *continuation*. When invoked,

Authors' Contact Information: Paulo Emílio de Vilhena, p.de-vilhena@imperial.ac.uk, Imperial College London, London, United Kingdom; Simcha van Collem, simcha.van.collem@ru.nl, Radboud University Nijmegen, Nijmegen, The Netherlands; Ines Wright, ines.w@cs.au.dk, Aarhus University, Aarhus, Denmark; Robbert Krebbers, mail@robbertkrebbers.nl, Radboud University Nijmegen, Nijmegen, The Netherlands.



This work is licensed under a Creative Commons Attribution 4.0 International License.

© 2026 Copyright held by the owner/author(s).

ACM 2475-1421/2026/1-ART34

<https://doi.org/10.1145/3776676>

the continuation resumes the suspended program, but, as a first-class value, the continuation can also be discarded or stored in memory to be invoked later.

The ability to suspend and resume programs can be used to implement interesting features such as coroutines [13] and promise-style asynchronous-programming libraries [19]. However, the ability to manipulate continuations is also dangerous. A continuation can capture resources. It may also contain code that must eventually be called to free up these resources. So, if the continuation is discarded, if it becomes unreachable, or if, for some other reason, it is not invoked, then some resources may never be released. Users of effect handlers must also make sure that the operation of performing an effect is always enclosed by a handler, otherwise, like an uncaught exception, an *unhandled effect* would cause a runtime error. For these reasons, effect handlers are widely seen as an advanced programming feature to be used with care.

An impressive range of tools to help programmers to reason about programs with effect handlers and to avoid these programming errors has been introduced. Programming languages such as Koka [37], Links [12, 26] and Effekt [10], for example, have type systems that statically ensure *effect safety*: unhandled effects are statically ruled out. Multiple other type systems with similar guarantees, covering a comprehensive portion of the design space of handlers, can be found in the literature [4–7, 11, 16, 32, 39, 51, 52, 57, 60].

In this paper, we are interested in expressing and verifying *relational properties* of programs with handlers, namely *program refinement* and *program equivalence*. These relational properties have several interesting applications. One could specify a complex but efficient algorithm or data structure in terms of a simple but inefficient counterpart, or express the correctness condition of *linearizability* for concurrent programs using program refinement [21]. Relational reasoning also plays a key role in compiler verification [2, 24].

The study of relational properties of programs with effect handlers is not new. Building on a logic to reason about equality of programs using effects described by an *algebraic theory* [44], Plotkin and Pretnar [46] introduce the notion of correctness of an effect handler as a relational property: the handler implementation must validate the equations of the corresponding algebraic theory. This seminal work has spawned a fertile investigation of relational logics for effect handlers [40, 49].

In prior work, relational reasoning is limited to a strictly functional setting deprived of built-in imperative features. To verify the correctness of a handler implementation that makes use of imperative features such as mutable state, the algebraic approach of Plotkin and Pretnar [46] requires the user to parameterize the correctness statement with an algebraic theory of state. Although denotational models for *ground store* [31] (that is, store in which cells can hold integers, pairs, sums, and references to other cells, but not functions or continuations) exist, to our knowledge unrestricted higher-order store still lacks an algebraic treatment. This limitation precludes the application of previous relational-reasoning approaches to programming languages like Links and OCaml, which, in addition to user-defined effects and handlers, have ready support for heap-allocated mutable state. The ability to store continuations on the heap is crucial in the effect-handler-based implementations of asynchronous-programming libraries that we study in this paper (§5.1). Moreover, although user-defined effects and handlers offer a modular basis for effectful programming, it is often the case that the handler-based implementation of an effect is obscured by the combination of advanced programming patterns, whereas its handler-free implementation can be derived directly using imperative features. Therefore, from a reasoning perspective, it is desirable to establish a formal statement relating a user-defined effect to its imperative counterpart. For example, can the operation **perform Fork** task, which performs the user-defined effect **Fork**, be seen as the operation **fork** (task()), which directly spawns a new thread?

To overcome these limitations and address this question, we introduce *blaze*, the first relational logic for a language supporting effect handlers, heap-allocated state, and primitive concurrency

and also the first *relational separation logic* for effect handlers. We build blaze on top of the Iris framework [28–30, 33–35] in the Rocq prover [54], thus providing users with the comfort of a proof assistant, the confidence of a mechanised theory, and the expressiveness of Iris, a modern higher-order concurrent separation logic with powerful features such as support for higher-order functions, user-defined ghost state, and invariants.

The blaze logic in a nutshell. The refinement relation $e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\}$ of blaze informally states that either e_l diverges, or both e_l and e_r terminate with values v_l and v_r that satisfy the postcondition R . The key ingredient is the parameter \mathcal{T} , which specifies the *relational theory* under which the refinement holds. This notion is inspired by Biernacki et al. [5], but, whereas they use pure (step-indexed [1, 3, 20]) logic to express relational theories, we use separation logic [42, 48], and, whereas they reason at the level of a *transparent* logical interpretation of types, we rely on abstract reasoning rules to manipulate an *opaque* notion of refinement. The novelty of blaze therefore *does not* lie in the construction of its model, which follows Biernacki et al. [5], but on the design of reasoning rules that allow the relational verification of handlers at a high level of abstraction hiding any model-specific details from the user of the logic. Moreover, by building this logic on top of separation logic, we are able to express relational properties that involve primitive effects of the language and that are conditional on the ownership of locations in the heap or on Iris-style ghost state.¹

Using relational theories we can relate user-defined effects to other user-defined effects, or relate user-defined effects to the native imperative features of the language. Concrete examples include:

- (1) Relating the *state effect* to the composition of *reader* and *writer effects*.
- (2) Relating the *state effect* effect to primitive load and store operations (§2).
- (3) Relating a handler-based implementation of *concurrency* to *true concurrency* (§5.1).
- (4) Expressing *algebraic laws*, for example that the non-deterministic choice operator (either implemented using a handler that collects a list of results or implemented using concurrency) satisfies monoid laws (§5.2).

Biernacki et al. [5, §4.2] already support (1). We port their result to blaze as part of our Rocq formalisation [17]. More crucially, by using separation logic to formulate our relational theories, blaze also supports (2) and (3). Another application of blaze is (4), which enables the formulation of a handler-correctness criterion in the style of Plotkin and Pretnar [46]. Expressing handler correctness in this style is novel in the context of higher-order state and primitive concurrency. However, unlike Plotkin and Pretnar [46]’s algebraic theories, we note that algebraic theories expressed in blaze are not transitive due to a known limitation of step-indexed relational logics [8, 27].

We give a semantics to the refinement relation $e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\}$ using an interpretation in Iris. At the basis, we use Iris’s weakest precondition assertion to define *observational refinement* in the same way as ReLoC [23, §7.1]. Then, taking inspiration from Pitts and Stark [43]’s *biorthogonality* technique (used for the first time by Biernacki et al. [5] in the context of effect handlers) we define refinement, mutually inductively with two other relations, using Iris’s guarded fixpoint operator.

While this layering of definitions makes it possible to bootstrap blaze, it also makes it infeasible to carry out refinement proofs directly by unfolding these definitions, let alone carry out these proofs in a compositional manner. We therefore take inspiration from ReLoC [22, 23] and Simuliris [2, 24] to develop a *relational logic* with a range of high-level reasoning principles that abstract over the details of these definitions. Our high-level logic provides a number of novel features:

¹Adding support for primitive effects, particularly concurrency, in a relational separation logic is not as straightforward as it may sound. Our proof rules for state directly follow ReLoC [23]. However, as we discuss in §5.1.1, such an approach does not go as smoothly for concurrency. We instead design original rules for reasoning about concurrency.

- (1) Our novel *introduction* and *exhaustion rules* make it possible to abstractly manipulate a relational theory \mathcal{T} . If \mathcal{T} contains a relation between e_l and e_r , then the introduction rule allows us to prove $e_l \lesssim e_r, \langle \mathcal{T} \rangle \{R\}$. The exhaustion rule allows us to eliminate the dependency on a theory \mathcal{T} , provided that the relations included in \mathcal{T} are correctly handled.
- (2) The *bind rule* makes it possible to focus on a subexpression and then continue with the verification of the whole expression in which the subexpression is replaced with a value. It is well known that, in the context of logics for effect handlers, a restriction on the bind rule is necessary for soundness. We develop a new restriction that requires the bound contexts to be *traversable* with respect to the relational theory \mathcal{T} . This flexibility is crucial to support dynamic effects labels.

We show the versatility of our approach through various extensions. We add support for dynamic labels in the style of OCaml’s `let exception` construct, following de Vilhena and Pottier [16]. Moreover, we add support for both one-shot and multi-shot continuations, taking inspiration from van Rooij and Krebbers [57]. Finally, as a side contribution needed to carry out some of our case studies, we introduce new relational rules for concurrency. These rules are independent of effect handlers and hold in any Iris-style relational logic such as ReLoC [22, 23].

Contributions. In sum, our contributions are the following:

- (1) **Novel relational logic.** We introduce blaze, the first relational separation logic for handlers.
- (2) **Case studies.** We conduct several challenging case studies including the verification that multiple effect-handler-based implementations of concurrency refine *truly concurrent* ones.
- (3) **Novel reasoning rules.** Our case studies led us to discover novel, simple, and general reasoning rules in relational concurrent separation logic that are independent of handlers.
- (4) **Correctness with respect to algebraic theories.** We show how the correctness of an effect handler with respect to an algebraic theory can be stated and proved in blaze.
- (5) **Mechanised theory.** We mechanise all our results, including soundness, in the Rocq prover.

2 Overview

In this section, we discuss the main challenges in designing a relational separation logic with support for effect handlers. Our goal is to informally explain how blaze handles these challenges. The examples are written in λ -blaze, a calculus whose syntax and semantics we explain in §3. In this section, we assume familiarity with functional programming and effect handlers. For the unaccustomed reader, Pretnar [47] provides a tutorial introduction to effect handlers.

Let us start by considering the following example:

```
countdown ≡ fun timer.
  timer.set 10; while (timer.get() > 0) {timer.set(timer.get() - 1)}
```

The function `countdown` receives an object `timer` as an argument with two fields, `get` and `set`. It assumes these fields implement the functionality to respectively access and update `timer`’s memory. It uses this functionality to update the timer from 10 to 0 through decrements of 1.

The definition of `countdown` is modular on the implementation of the timer. In a language with effect handlers, the programmer can exploit this generality by implementing `get` and `set` as user-defined effects and providing different handlers to customise the implementation of the effects performed by `get` and `set`. For example, assuming an effect `$Timer` is available, a generic implementation of `get` and `set` can be obtained as follows:

```
timer ≡ {get = fun _ . perform $Timer (inl()); set = fun y . perform $Timer (inr y)}
```

The operations `get` and `set` perform the effect `$Timer`. The *left and right injections* `inl` and `inr` are used to distinguish requests sent by `get` and `set`. A `$Timer` handler eventually assigns meaning to `get` and `set` by replying to these requests. Here are two possible instances of such handlers:

```
run_st_passing ≡ fun main.
  let run = fun () .
    handle main() with
      | effect $Timer request, k ⇒ fun x .
        match request with
          | inl () ⇒ k x x
          | inr y ⇒ k () y
          | y ⇒ fun _ . y
    in run() 0

run_heap ≡ fun main.
  let r = ref 0 in
  handle main() with
    | effect $Timer request, k ⇒
      match request with
        | inl () ⇒ k (!r)
        | inr y ⇒ r ← y; k ()
        | y ⇒ y
  in run() 0
```

Both receive a piece of client code `main` that performs `$Timer` effects. The function `run_st_passing` installs a handler that interprets `$Timer` effects in *state-passing style*, whereby the computation is transformed into a function that takes the current state of the timer and outputs the timer's final state. In contrast, the function `run_heap` interprets `$Timer` effects by storing the current state of the timer in a local reference `r`. This implementation is arguably simpler than the state-passing implementation of `run_st_passing` although presumably they implement the same functionality.

This observation motivates a key question: is it possible to show that `run_st_passing` is a *refinement* of `run_heap`? That is, can `run_heap` be used as a *specification* of `run_st_passing` and, therefore, as a *reference implementation* of `get` and `set`?

The notion of refinement is formalised in relational logics as the relation $e_l \lesssim e_r \{R\}$, where e_l and e_r are expressions and the *postcondition* R is a relation on values. We refer to e_l as the expression on the *implementation side* and to e_r as the expression on the *specification side*. The refinement relation informally states that either e_l diverges or both e_l and e_r terminate with outputs v_l and v_r such that $R(v_l, v_r)$ holds, capturing the intuition that e_l implements the same functionality described by e_r , because, informally, every output of e_l corresponds to an output of e_r related by R .

The question can thus be reformulated as how to establish a refinement between `run_st_passing` and `run_heap`, such as the statement

$$\text{impl}_1 \lesssim \text{impl}_2 \{y_l \ y_r. \ y_l = y_r\}, \quad \text{where } \text{impl}_1 \triangleq \text{run_st_passing} (\text{fun}(). \ \text{countdown timer}) \quad \text{and } \text{impl}_2 \triangleq \text{run_heap} (\text{fun}(). \ \text{countdown timer}), \quad (1)$$

expressing the property that impl_1 and impl_2 have the same outputs.

To our knowledge, there are no relational logics with support for effect handlers *and* heap-allocated mutable state and therefore no logics where such a relation can be derived. Addressing this gap, we introduce blaze, the first relational separation logic with support for handlers. The choice of a separation logic enables modular reasoning about state-manipulating programs such as `run_heap`. The following subsections explain other interesting and novel aspects of blaze.

2.1 Modular Reasoning About Effects: Handler Versus Handlee

In blaze, it is possible to state and prove Refinement 1. In fact, this refinement can be established in a *compositional* way, whereby the proof is split into two parts: a proof that the handlers installed by `run_st_passing` and `run_heap` are related and a proof that the handlees monitored by these handlers are related. In the current example, this creates the following two subgoals:

$$\text{countdown timer} \stackrel{?}{\lesssim} \text{countdown timer} \quad (2)$$

$$\forall \text{main}_l, \text{main}_r. \ \text{main}_l() \stackrel{?}{\lesssim} \text{main}_r() \dashv \text{run_st_passing} \text{main}_l \lesssim \text{run_heap} \text{main}_r \{=\} \quad (3)$$

Refinement 2 relates the handlees and Refinement 3 relates *run_st_passing* and *run_heap* under the assumption they receive related arguments. For brevity, we write “=” in Refinement 3 for the postcondition $y_l \ y_r. \ y_l = y_r$. Moreover, we use $\overset{?}{\lesssim}$ to denote a notion of refinement that is yet to be defined. Recall that the informal reading of *countdown timer* $\lesssim _ \{ \}$ states *countdown timer* either diverges or terminates with a value. This standard notion of refinement $_ \lesssim _ \{ \}$ is therefore insufficient, because, without a handler, the program *countdown timer* performs unhandled effects.

This limitation reveals a *key challenge*: to reason about the handlee independently of the handler, it is necessary to generalise the standard notion of refinement to account for unhandled effects.

The blaze logic solves this challenge by parameterising the refinement relation with a *relational theory*. A relational theory can be seen as a set of assumed refinements. Concretely, it is defined as a set of triples (e_l, e_r, Q) , where e_l and e_r are expressions and Q is a relation on pairs of expressions called the *return condition*. In short, the return condition describes the condition under which e_l and e_r can return. For example, the return condition $y_l \ y_r. \ y_l = y_r$ states e_l and e_r can return only when they terminate with the same values. In this case, the return condition Q can be seen as a postcondition. The reading of (e_l, e_r, Q) then simply states that e_l refines e_r with postcondition Q . For the purposes of this section, this first approximation is enough.

The general refinement relation in blaze has the form $e_l \overset{?}{\lesssim} e_r \langle \mathcal{T} \rangle \{R\}$, where \mathcal{T} is the parameterised relational theory. When a relational theory is empty, we write $e_l \lesssim e_r \{R\}$ which has the same informal meaning as before. The informal reading of the general relation $e_l \overset{?}{\lesssim} e_r \langle \mathcal{T} \rangle \{R\}$ is that $K_l[e_l] \overset{?}{\lesssim} K_r[e_r] \{R\}$ holds for every pair of contexts K_l and K_r that validate the theory \mathcal{T} . A pair of contexts K_l and K_r validate \mathcal{T} when the refinements included in \mathcal{T} hold under K_l and K_r ; that is, if \mathcal{T} includes the relation between two expressions e_l and e_r , then $K_l[e_l]$ refines $K_r[e_r]$. This general notion of refinement allows us to reason about programs e_l and e_r that perform unhandled effects, because, when \mathcal{T} is well-chosen, the contexts K_l and K_r that validate \mathcal{T} are precisely those that handle the effects performed by e_l and e_r . At the same time, the contexts K_l and K_r appear only in the definition of $e_l \overset{?}{\lesssim} e_r \langle \mathcal{T} \rangle \{R\}$, which, during a verification task, need not be unfolded. The theory \mathcal{T} can thus be seen as a logical abstraction of the contexts under which e_l and e_r occur.

In the running example of Refinements 2 and 3, it is now possible to substitute $\overset{?}{\lesssim}$ with a refinement relation parameterised by a relational theory, say $Timer_{refl}$:

$$_ \overset{?}{\lesssim} _ \quad \triangleq \quad _ \lesssim _ \langle Timer_{refl} \rangle \{=\}$$

There are two minimal requirements for the relational theory $Timer_{refl}$: (1) $Timer_{refl}$ must include sufficiently many relations so that $countdown \ timer \overset{?}{\lesssim} countdown \ timer \langle Timer_{refl} \rangle \{=\}$ holds and (2) $Timer_{refl}$ must be sufficiently small so that $run_st_passing \ main_l \overset{?}{\lesssim} run_heap \ main_r \{=\}$ holds under the assumption that $main_l() \overset{?}{\lesssim} main_r() \langle Timer_{refl} \rangle \{=\}$ holds. A choice of $Timer_{refl}$ that satisfies both requirements is one that includes only the following refinement:

$$\forall v. \mathbf{perform} \$\text{Timer } v \overset{?}{\lesssim} \mathbf{perform} \$\text{Timer } v \langle Timer_{refl} \rangle \{=\} \quad (4)$$

This is sufficient to prove $countdown \ timer \overset{?}{\lesssim} countdown \ timer \langle Timer_{refl} \rangle \{=\}$, because, as the two expressions in this relation are same, every $\$ \text{Timer}$ effect on one side corresponds to exactly one $\$ \text{Timer}$ effect on the other side. Therefore, when reasoning about performing an effect, it suffices to apply Refinement 4 to conclude that, in both expressions, the results are the same.

Moreover, because $Timer_{refl}$ includes only Refinement 4, it follows that, if $e_l \overset{?}{\lesssim} e_r \langle Timer_{refl} \rangle \{=\}$ holds for arbitrary expressions e_l and e_r , then it must be the case that every $\$ \text{Timer}$ effect in e_l corresponds to exactly one $\$ \text{Timer}$ effect in e_r . This assumption can be exploited by the proof of $run_st_passing \ main_l \overset{?}{\lesssim} run_heap \ main_r \{=\}$ to establish the relation between the two handlers.

2.2 Flexible Reasoning: Handler-Based Versus Handler-Free Implementations

One of the motivations to establish the refinement between *run_st_passing* and *run_heap* is that *run_heap* provides a simpler and more direct implementation of the timer when compared to the state-passing implementation of *run_st_passing*. However, *run_heap* does not exploit non-trivial functionalities of effect handlers as the effect branch always immediately resumes the continuation. This observation permits an implementation of the timer without effects:

$$\text{ref_timer} \triangleq \mathbf{fun} r. \{ \text{get} = \mathbf{fun} () . !r; \text{set} = \mathbf{fun} y. r \leftarrow y \}$$

The question now is: can the refinement

$$\text{impl}_1 \lesssim \text{impl}_3 \{=}, \text{ where } \text{impl}_3 \triangleq \mathbf{let} r = \mathbf{ref} 0 \mathbf{in} \text{countdown}(\text{ref_timer} r), \quad (5)$$

be established in blaze? Moreover, if possible, can the proof be done in a compositional way like in the previous example, where reasoning about handlee and handler are carried out independently?

The answers to both questions are positive: the refinement can be established in blaze with a compositional proof. Indeed, the proof of $\text{impl}_1 \lesssim \text{impl}_3 \{=\}$ is split into two subgoals:

$$\forall \ell. \ell^{1/2} \mapsto_s 0 \rightarrow * \text{countdown timer} \lesssim \text{countdown}(\text{ref_timer} \ell) \langle \text{Timer}_{\text{spec}}^{\ell} \rangle \{=\} \quad (6)$$

$$\forall \ell, \text{main}_l, e_r. \ell^{1/2} \mapsto_s 0 \rightarrow * \text{main}_l() \lesssim e_r \langle \text{Timer}_{\text{spec}}^{\ell} \rangle \{=\} \rightarrow * \text{run_st_passing main}_l \lesssim e_r \{=\} \quad (7)$$

Refinement 6 relates the handlee *countdown timer* to *countdown*(*ref_timer* ℓ) under the theory $\text{Timer}_{\text{spec}}^{\ell}$, which we introduce shortly. Refinement 7 is stated in an interesting way. It relates *run_st_passing main_l* to an arbitrary expression e_r . Intuitively, the expression represents the program *countdown*(*ref_timer* ℓ), but, thanks to the theory $\text{Timer}_{\text{spec}}^{\ell}$, this specific program can be entirely abstracted: all the information needed to carry out Refinement 7 is that *main_l* refines e_r under $\text{Timer}_{\text{spec}}^{\ell}$.

The variable ℓ stands for the location to which r (in impl_3) is bound. As usual in relational separation logics, each of the two programs in a refinement relation manipulates its own heap. The *points-to predicate* $_ \mapsto_s _$ describes the state of the heap of the program on the specification side, whereas $_ \mapsto_i _$ describes the state of the heap on the implementation side. The fraction that appears on top of \mapsto_s represents a *fractional ownership* [9] of ℓ : it grants read-only permission to ℓ . Full ownership can be retrieved by combining two $\ell^{1/2} \mapsto_s _$ assertions. In the proof of Refinement 5, fractional assertions $\ell^{1/2} \mapsto_s _$ are given to both the handlee and the handler. Full ownership is therefore retrieved when the handlee performs an effect and ownership of the handlee's fractional assertion $\ell^{1/2} \mapsto_s _$ is temporarily transferred to the handler until the handlee is resumed.

Like $\text{Timer}_{\text{refl}}$, the theory $\text{Timer}_{\text{spec}}^{\ell}$ must fulfil two requirements: (1) the theory must be sufficiently relaxed so that Refinement 6 can be established and (2) it must be sufficiently small so that the terms $\text{main}_l()$ and e_r in Refinement 7 are tightly related. The first requirement now seems particularly challenging because Refinement 6 relates an effectful program to a non-effectful one. Fortunately, relational theories are not limited to relations between only effectful expressions like in $\text{Timer}_{\text{refl}}$. They can in fact express relations between arbitrary expressions. Taking advantage of this flexibility, the theory $\text{Timer}_{\text{spec}}^{\ell}$ includes a relation between the effectful implementation of *get* and *set* fields of *timer* and their heap-manipulating counterparts of *ref_timer*:

$$\forall x. \ell^{1/2} \mapsto_s x \rightarrow * \mathbf{perform} \$\text{Timer} (\mathbf{inl} ()) \lesssim !\ell \langle \text{Timer}_{\text{spec}}^{\ell} \rangle \{ y_l y_r. y_l = y_r = x * \ell^{1/2} \mapsto_s x \} \quad (8)$$

$$\forall y. \ell^{1/2} \mapsto_s _ \rightarrow * \mathbf{perform} \$\text{Timer} (\mathbf{inr} y) \lesssim \ell \leftarrow y \langle \text{Timer}_{\text{spec}}^{\ell} \rangle \{ y_l y_r. y_l = y_r = () * \ell^{1/2} \mapsto_s y \} \quad (9)$$

From the perspective of the handlee, these relations guarantee that performing the effect $\$T$ imer is similar to manipulating the memory location ℓ .

2.3 Context-Local Relational Reasoning

A closer look at Refinements 8 and 9 reveals an important limitation. They apply only to pairs of programs e_l and e_r where e_l consists precisely of a single $\$Timer$ effect and e_r consists precisely of a single read or store operation. As such, they are insufficient to establish Refinement 6, because the calls to get and set in *countdown timer* and in *countdown (ref_timer ℓ)* occur in the context of a larger program, not as single operations.

The key missing principle to address this limitation is the *bind rule*. The bind rule allows the user to reason about a piece of code independently of the context under which this code is eventually executed. In standard relational logics, the bind rule is formally stated as follows:

$$\text{STANDARD-BIND} \quad e_l \lesssim e_r \{y_l y_r. K_l[y_l] \lesssim K_r[y_r] \{R\}\} \vdash K_l[e_l] \lesssim K_r[e_r] \{R\}$$

This rule is sound in blaze, but insufficient because it assumes the parameterised theory is empty. A natural fix would be to decorate every occurrence of the refinement relation with a theory \mathcal{T} :

$$\text{UNSOUND-BIND} \quad e_l \lesssim e_r \langle \mathcal{T} \rangle \{y_l y_r. K_l[y_l] \lesssim K_r[y_r] \langle \mathcal{T} \rangle \{R\}\} \vdash K_l[e_l] \lesssim K_r[e_r] \langle \mathcal{T} \rangle \{R\}$$

The resulting rule is *unsound*. To see why, it suffices to consider the following counterexample, where the effect $\$Id$ is assumed to be available:

$$e_{true} \triangleq \text{handle } (\text{perform } \$Id \text{ true}) \text{ with effect } \$Id x, k \Rightarrow k x | y \Rightarrow y$$

If we further assume there is a theory Neq that includes the refinement $\forall b \in \text{Bool}. \text{perform } \$Id b \lesssim \text{perform } \$Id b \langle Neq \rangle \{\neq\}$, then, using UNSOUND-BIND with both K_l and K_r instantiated as the $\$Id$ handler, it is possible to establish the refinement $e_{true} \lesssim e_{true} \langle Neq \rangle \{\neq\}$, which is false, because both sides of the refinement terminate with true.

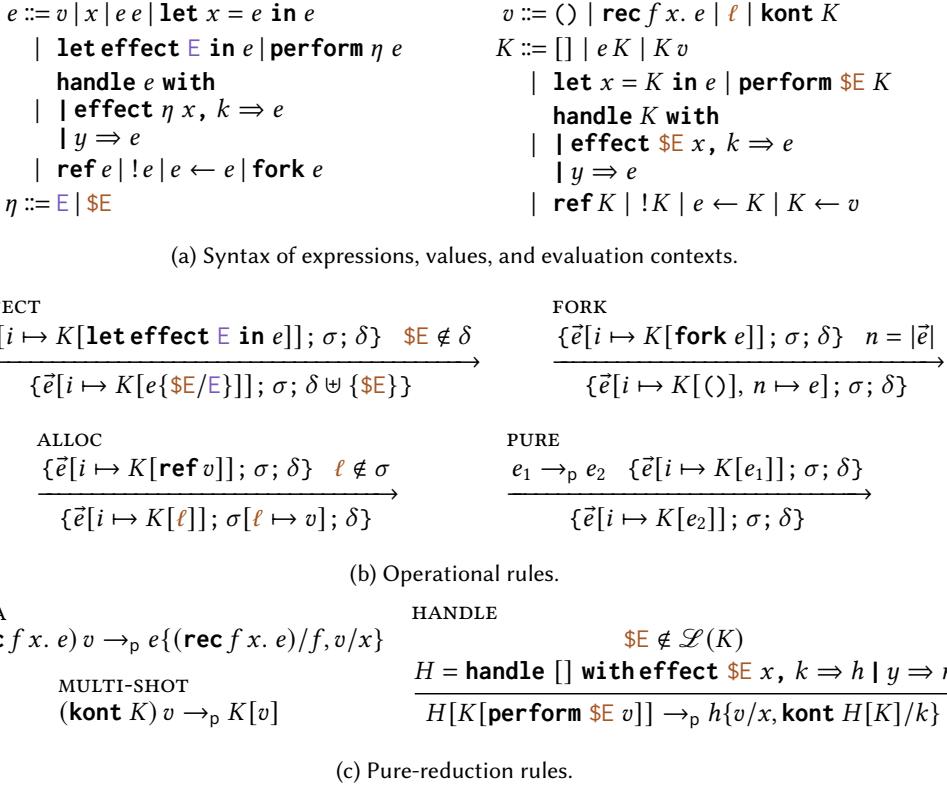
This counterexample suggests that to enable sound context-local reasoning, there must be some restriction on the evaluation contexts K_l and K_r . In particular, the rule should not be applicable when the contexts K_l and K_r contain handlers for the effects described by the theory \mathcal{T} . In blaze, a sound bind rule integrating these restrictions is formulated as follows:

$$\text{BIND} \quad \begin{array}{c} \text{traversable}(K_l, K_r, \mathcal{T}) \\ e_l \lesssim e_r \langle \mathcal{T} \rangle \{y_l y_r. K_l[y_l] \lesssim K_r[y_r] \langle \mathcal{F} \rangle \{R\}\} \end{array} \vdash K_l[e_l] \lesssim K_r[e_r] \langle \mathcal{F} \rangle \{R\}$$

The rule is applicable when there exists a theory \mathcal{T} *included* in \mathcal{F} , such that $\text{traversable}(K_l, K_r, \mathcal{T})$ holds. The predicate $\text{traversable}(K_l, K_r, \mathcal{T})$ intuitively states that K_l and K_r do not conflict with \mathcal{T} , or, visually, that \mathcal{T} can *traverse* K_l and K_r . It is defined in an abstract way with no reference to the handlers in K_l and K_r . In the case of Timer_{refl} , it is possible to show this predicate holds for any contexts K_l and K_r that contain no $\$Timer$ handler. In the case of Timer_{spec}^ℓ , the predicate holds for any K_l that contains no $\$Timer$ handler. No condition is imposed on K_r in this case, because the expressions on the right-hand side of Refinements 8 and 9 do not perform effects. The blaze logic therefore enjoys a powerful context-local reasoning principle that is adjustable to the parameterised theory. As we are going to show in §4.2, this principle is especially important to support reasoning in the presence of multiple effect names.

3 Language

We introduce λ -blaze, an untyped calculus with formally defined syntax and semantics. The language has support for heap-allocated mutable state and concurrency, both deep and shallow handlers, both one-shot and multi-shot continuations, and dynamically allocated effect names. For most of the paper, only deep handlers that capture multi-shot continuations are used. So, for the sake of conciseness, we postpone the introduction of the syntax and semantics of one-shot continuations to §4.3, where we explain the extension of the logic with support for this feature.

Fig. 1. Syntax and semantics of λ -blaze.

3.1 Syntax

Figure 1a shows the syntax of expressions, values, and evaluation contexts. The definition of evaluation contexts reflects a right-to-left evaluation order. Every node in the syntax tree of an evaluation context K contains exactly one child, except for the empty context which contains none. Thanks to this observation, a context K can be seen as a list whose elements, called *frames*, are the nodes in its syntax tree. We use the notation $K[K']$ to denote the context obtained by concatenating these lists. The similar notation $K[e]$ is used to denote the expression obtained by the operation of *filling* K with e , characterised by the equations: $[] [e] = e$ and $(K[K']) [e] = K[K'[e]]$.

Most of the syntactic constructs of the language are standard. In the following paragraphs, we explain two aspects that are unusual: the syntax of function definitions and the distinction between *effect names* and *effect labels*.

Function definitions. Functions are defined using the syntax $\text{rec } f x. e$. The variable x is a formal argument of the function. Its scope is e . The variable f is bound to the function definition itself (that is, the entire term $\text{rec } f x. e$). It can be used in the scope of e to write recursive definitions. When f does not occur in e , we use the simpler notation $\text{fun } x. e$. For function definitions with more than one formal argument, we introduce the following syntactic sugar: $\text{fun } \vec{x}. e \triangleq \text{fun } x_0. \dots \text{fun } x_{n-1}. e$ and $\text{rec } f \vec{x}. e \triangleq \text{rec } f x_0. \text{fun } \vec{x}_1 \dots x_{n-1}. e$, where $n = |\vec{x}|$.

Effect names and effect labels. Taking inspiration from previous work [7, 16, 60], the syntax of λ -blaze makes a clear distinction between effect names E and effect labels $\$E$. Effect names are binders whose scope is delimited by the construct **let effect** E **in** e (following a syntax similar to OCaml’s **let exception** construct). Effect labels appear at runtime after the execution of a **let effect** construct, which binds an effect name to an effect label. The motivation for introducing this distinction is to provide λ -blaze with mechanisms to avoid the issue of *colliding effect names* [7, 16, 60], when the same effect name is used in two unrelated pieces of code.

To give an example, consider the following implementation of an *ask effect* [5] and the client code *colliding*, which installs an **Ask** handler over calls to the function it receives as an argument:

```
run_ask ≡ fun x main. let effect Ask in let ask = fun _ . perform Ask () in
    handle main ask with effect Ask (), k ⇒ k x | z ⇒ z
colliding ≡ fun ask₀. let effect Ask in let ask₁ = fun _ . perform Ask () in
    handle ask₀ () + ask₁ () with effect Ask (), k ⇒ k 1 | z ⇒ z
```

The collision occurs during the execution of *run_ask* 0 *colliding*. The call to *ask*₀ in *colliding* performs an **Ask** effect that should be handled by *run_ask*, but, at this moment, the innermost handler is the one installed by *colliding*. If effect names were used to find the handlers, then the call to *ask*₀ would be handled by *colliding*’s handler.

This example illustrates the issue of collision of effect names: the name **Ask** is used with different purposes by two unrelated pieces of code. It would thus be desirable for the semantics of **let effect** E **in** e to take care of avoiding this collision of names. This is exactly what it does: when **let effect** E **in** e is executed, it allocates a fresh *effect label* $\$E$ which is substituted for E in e . At runtime, it is $\$E$ that is used to perform and handle effects. According to this semantics, the program *run_ask* 0 *colliding* runs as expected, because, at runtime, the handlers installed by *run_ask* and *colliding* handle effects for different labels.

3.2 Semantics

The semantics is defined using three sorts of runtime terms: memory locations ℓ , created by **ref** instructions; multi-shot continuations **kont** K , created by handlers during the handling of an effect; and effect labels $\$E$, which, as explained, are created by **let effect** instructions.

Memory operations follow a standard heap semantics. The semantics of handlers that capture multi-shot continuations is also standard [47]. The semantics of **let effect** follows de Vilhena and Pottier [16]. The semantics is formalised by a number of operational rules, of which the most relevant can be seen in Figure 1b. The rules define a reduction relation between *configurations* of the form $\{\vec{e}; \sigma; \delta\}$, where \vec{e} is a pool of running threads, σ is a store, and δ is a set of allocated effect labels. Rule **ALLOC** captures the semantics of memory allocation where a fresh location ℓ is non-deterministically chosen and initialised in σ with v . Rule **FORK** captures the semantics of **fork** e , allocates a thread to execute e and returns $()$. Rule **EFFECT** captures the semantics of **let effect**: to guarantee freshness, the non-deterministically chosen label $\$E$ must not be in the set of pre-allocated labels δ . Finally, Rule **PURE** captures the semantics of *pure reductions* $e \rightarrow_p e'$, partially defined in Figure 1c. Rule **BETA** is the standard beta reduction. Rule **MULTI-SHOT** shows how the invocation of a multi-shot continuation **kont** K restores K as an evaluation context. Rule **HANDLE** shows how control is transferred to the effect branch h of a handler in case of an effect. The term $\mathcal{L}(K)$ denotes the labels of the handlers in K . The condition $\$E \notin \mathcal{L}(K)$ ensures H is the innermost handler.

4 Logic

Our logic consists of two main layers with independent notions of refinement and independent reasoning rules. The first layer, baze, offers a base logic built directly on top of Iris [29]. The

baze logic offers an expressive notion of refinement for arbitrary λ -blaze programs. Reasoning about programs with multiple effect labels in baze however can be challenging, motivating the introduction of the second layer, blaze, which is built on top of baze and tailored for programs with dynamic labels. So far, we have used the name blaze for the collection of both logics. To avoid confusion, from now on, we use the term blaze to refer exclusively to this second layer.

As noted in §1, the novelty of both logical layers lies in the design of a comprehensive set of high-level reasoning rules (Figures 3 and 5) for the relational verification of programs with handlers. The model of baze closely follows Biernacki et al. [5]’s model of refinement, Allain et al. [2]’s domain of *abstract protocols*, and ReLoC [23]’s model of *observational refinement*. Moreover, the model of blaze closely follows the model of TesLogic [16].

4.1 baze: The Base Logic

The refinement statement in baze takes the form $e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\}$, where e_l and e_r are λ -blaze programs, \mathcal{T} is a parameterised relational theory, and the postcondition R is a predicate on pairs of values. It intuitively means that, under any pair of contexts K_l and K_r that validate the theory \mathcal{T} , either $K_l[e_l]$ diverges or both $K_l[e_l]$ and $K_r[e_r]$ terminate with values v_l and v_r such that $R(v_l, v_r)$ holds. The key to formalise this notion of refinement is thus to precisely formulate what is a theory \mathcal{T} and what it means for a pair of contexts to validate \mathcal{T} .

4.1.1 Relational Theories. A theory \mathcal{T} is modelled as a predicate of type²

$$iThy \triangleq (Expr \times Expr \times (Expr \times Expr \rightarrow iProp)) \rightarrow iProp,$$

where $iProp$ is the type of Iris assertions.³ Intuitively, the assertion $\mathcal{T}(e_l, e_r, Q)$ means that e_l is related to e_r and that e_l and e_r can be replaced with any pair of expressions e'_l and e'_r for which the return condition $Q(e'_l, e'_r)$ holds. Perhaps the simplest example of a relational theory is the *empty theory* \perp , which does not include any relations: $\perp(e_l, e_r, Q) \triangleq \text{False}$.

For a slightly more involved example, consider the definition of theory $Timer_{refl}$ from §2.1 relating the effect $\$Timer$ ⁴ to itself and asserting that both effects return the same output:⁵

$$Timer_{refl}(\mathbf{perform } \$Timer v, \mathbf{perform } \$Timer v, Q) \triangleq \square \forall w \in Val. Q(w, w)$$

To express that both $\mathbf{perform } \$Timer v$ operations return the same output, the theory $Timer_{refl}$ asserts the return condition Q holds of any pair of copies of the same value: $\square \forall w \in Val. Q(w, w)$. This assertion is guarded by Iris’s *persistently modality* \square . Typical separation-logic assertions, such as the points-to assertion $\ell \mapsto v$, declare ownership of resources, so, by default, they cannot be arbitrarily shared or duplicated. The persistently modality indicates when an assertion does not claim ownership of ephemeral resources and thereby *can* be duplicated. Here, it is used to indicate that the effect $\$Timer$ complies with a multi-shot policy whereby the operation $\mathbf{perform } \$Timer v$ can return multiple times. The return condition must hold every time the operation returns.

4.1.2 Context-Closure Operation. As defined, the theory $Timer_{refl}$ suffers from the limitation highlighted in §2.3: $Timer_{refl}$ is limited to relations between single $\mathbf{perform }$ expressions, when, in fact, it is desirable for the theory to enjoy some form of context-local reasoning with which $\mathbf{perform }$ expressions can be related under evaluation contexts. More abstractly, we wish to *close a theory* \mathcal{T}

²This type is similar to the type of semantic rows \mathbf{Eff} from Biernacki et al. [5, §3.2] and to the type of *abstract protocols* from Allain et al. [2, §6]. See §6 for an in-depth discussion.

³Iris assertions include standard connectives and quantifiers, separation logic connectives (in particular, the *separating conjunction* $*$ and the *separating implication* $\dashv*$), and modalities whose purpose and meaning we explain as they appear.

⁴In §2.1, we assume $\$Timer$ is *available*. Formally, this assumption means $\$Timer$ is created by a **let effect** instruction placed at the global level.

⁵We pattern match on the expressions e_l and e_r in $\mathcal{T}(e_l, e_r, Q)$, yielding False if no pattern applies.

$$\begin{aligned}
O(e_l, e_r, S) &\triangleq \forall i, K. \text{specCtx} \multimap i \Rightarrow K[e_r] \multimap \text{wp } e_l \{v_l\} \exists v_r. i \Rightarrow K[v_r] * S(v_l, v_r) \\
e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\} &\triangleq \forall K_l, K_r, S. \{R\} K_l \lesssim K_r \langle \mathcal{T} \rangle \{S\} \multimap O(K_l[e_l], K_r[e_r], S) \\
\{R\} K_l \lesssim K_r \langle \mathcal{T} \rangle \{S\} &\triangleq \wedge \begin{cases} \forall v_l, v_r. R(v_l, v_r) \multimap O(K_l[v_l], K_r[v_r], S) \\ \forall e_l, e_r. \mathcal{T} \blacktriangleleft e_l \lesssim e_r \{R\} \multimap O(K_l[e_l], K_r[e_r], S) \end{cases} \\
\mathcal{T} \blacktriangleleft e_l \lesssim e_r \{R\} &\triangleq \exists Q. \mathcal{T}(e_l, e_r, Q) * \square \triangleright \forall e'_l, e'_r. Q(e'_l, e'_r) \multimap e'_l \lesssim e'_r \langle \mathcal{T} \rangle \{R\}
\end{aligned}$$

Fig. 2. Model of baze.

under contexts, so that, along rough lines, if $\mathcal{T}(e_l, e_r, Q)$ holds, then so does $\mathcal{T}(K_l[e_l], K_r[e_r], P)$ for some return condition P . To this end, we introduce the *context-closure* operation:

$$\begin{aligned}
((ls_l, ls_r) \Downarrow \mathcal{T})(e_l, e_r, P) &\triangleq \\
\exists e'_l, e'_r, K_l, K_r, Q. & \left(\begin{array}{l} e_l = K_l[e'_l] * e_r = K_r[e'_r] * \text{neutral}(ls_l, K_l) * \text{neutral}(ls_r, K_r) * \\ \mathcal{T}(e'_l, e'_r, Q) * \square \forall e''_l, e''_r. Q(e''_l, e''_r) \multimap P(K_l[e''_l], K_r[e''_r]) \end{array} \right)
\end{aligned}$$

The context-closure of a theory \mathcal{T} enables the relation of expressions of the form $K_l[e'_l]$ and $K_r[e'_r]$, provided the subexpressions e'_l and e'_r are related. A common restriction on the contexts K_l and K_r under which e'_l and e'_r appear is that they contain no handlers for the effects performed by these expressions. To incorporate this restriction, the context-closure operation is parameterised by a pair of lists of labels ls_l and ls_r and includes the condition that K_l and K_r be respectively *neutral* for ls_l and ls_r . A context K is neutral for a list of labels ls , noted $\text{neutral}(ls, K)$, when K contains no handlers for labels in ls : $\mathcal{L}(K) \cap ls = \emptyset$.

In the example of $\text{Timer}_{\text{refl}}$, the context-closure $\text{Timer}'_{\text{refl}} \triangleq ([\$Timer], [\$Timer]) \Downarrow \text{Timer}_{\text{refl}}$ enables the relation of expressions under contexts K_l and K_r respectively neutral for $\$Timer$: $\text{Timer}'_{\text{refl}}(e_l, e_r, \lambda e'_l e'_r. Q(K_l[e'_l], K_r[e'_r])) \vdash \text{Timer}'_{\text{refl}}(K_l[e_l], K_r[e_r], Q)$.

4.1.3 Model. With the definition of $i\text{Thy}$, it is now possible to formalise the notion of validation of a theory by a pair of contexts and consequently to formalise the notion of a refinement relation parameterised by a theory.

Figure 2 shows the definition of the refinement relation $e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\}$. The definition is recursive and relies on the validation of a theory by a pair of contexts $\{R\} K_l \lesssim K_r \langle \mathcal{T} \rangle \{S\}$. The definition also relies on the notions of *observational refinement* $O(e_l, e_r, S)$ and of *admissibility of a refinement by a theory* $\mathcal{T} \blacktriangleleft e_l \lesssim e_r \{R\}$. The notion of admissibility is transparent to the user, whereas the notions of observational refinement and theory validation are opaque and used only in the model.

Observational refinement is defined exactly like in ReLoC [23, §7.1] (which is inspired by Turon et al. [56]). The intuitive reading of $O(e_l, e_r, S)$ is that either e_l diverges or both e_l and e_r respectively terminate with values v_l and v_r such that $S(v_l, v_r)$ holds. The formal definition makes use of Iris's weakest precondition $\text{wp } e_l \{\dots\}$, which expresses precisely the condition that e_l either diverges or terminates with a value. To express conditions on e_r , the definition makes use of the *ghost thread-pool* assertion $i \Rightarrow K[e_r]$ to state that thread i on the specification side is about to execute e_r . The use of $i \Rightarrow K[v_r]$ as a postcondition means that thread i finished executing e_r and that e_r returned output v_r .⁶ The thread identifier i is universally quantified, because it is not particularly

⁶The assertion specCtx is used to momentarily give ownership of the specification side's resources. It is defined like in Frumin et al. [23]. Its definition is included in the Appendix [18, Figure 14].

VALUE $\frac{R(v_l, v_r)}{v_l \lesssim v_r \langle \mathcal{T} \rangle \{R\}}$	INTRODUCTION $\frac{\mathcal{T} \blacktriangleleft e_l \lesssim e_r \langle R \rangle}{e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\}}$	BIND $\frac{\begin{array}{c} traversable(K_l, K_r, \mathcal{T}) \\ \mathcal{T} \sqsubseteq \mathcal{F} \\ e_l \lesssim e_r \langle \mathcal{T} \rangle \{v_l v_r. K_l[v_l] \lesssim K_r[v_r] \langle \mathcal{F} \rangle \{R\}\} \end{array}}{K_l[e_l] \lesssim K_r[e_r] \langle \mathcal{F} \rangle \{R\}}$
EXHAUSTION $\frac{\begin{array}{c} e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\} \\ \wedge \left\{ \begin{array}{l} \forall v_l, v_r. R(v_l, v_r) \rightarrow K_l[v_l] \lesssim K_r[v_r] \langle \mathcal{F} \rangle \{S\} \\ \forall e'_l, e'_r. \mathcal{T} \blacktriangleleft e'_l \lesssim e'_r \langle R \rangle \rightarrow K_l[e'_l] \lesssim K_r[e'_r] \langle \mathcal{F} \rangle \{S\} \end{array} \right. \end{array}}{K_l[e_l] \lesssim K_r[e_r] \langle \mathcal{F} \rangle \{S\}}$		MONOTONICITY $\frac{\begin{array}{c} e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\} \\ \square \forall v_l, v_r. R(v_l, v_r) \rightarrow S(v_l, v_r) \end{array}}{e_l \lesssim e_r \langle \mathcal{F} \rangle \{S\}}$
STEP-L $\frac{e_l \rightarrow_p e'_l \quad \triangleright K[e'_l] \lesssim e_r \langle \mathcal{T} \rangle \{R\}}{K[e_l] \lesssim e_r \langle \mathcal{T} \rangle \{R\}}$	STEP-R $\frac{e_r \rightarrow_p e'_r \quad e_l \lesssim K[e'_r] \langle \mathcal{T} \rangle \{R\}}{e_l \lesssim K[e_r] \langle \mathcal{T} \rangle \{R\}}$	

Fig. 3. Reasoning rules of the base logic.

relevant which specific thread is related to e_l as long as it executes e_r . The context K under which e_r runs is universally quantified to endow observational refinement with context-local reasoning.⁷

Given the intuitive reading of observational refinement $\mathcal{O}(e_l, e_r, S)$ and assuming that the notion of validation of a theory \mathcal{T} by a pair of contexts K_l and K_r is captured by $\{R\} K_l \lesssim K_r \langle \mathcal{T} \rangle \{S\}$, the definition of the refinement relation $e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\}$ reads as: for every pair of contexts K_l and K_r , and for all postconditions S , if K_l and K_r validate \mathcal{T} , then, either $K_l[e_l]$ diverges or both $K_l[e_l]$ and $K_r[e_r]$ terminate with values by S . The universal quantification over contexts is inspired by the *biorthogonality* technique of Pitts and Stark [43], used for the first time by Biernacki et al. [5] to define logical relations for a language with effect handlers.

The definition of $\{R\} K_l \lesssim K_r \langle \mathcal{T} \rangle \{S\}$ consists of the conjunction of two clauses: (1) a clause relating K_l and K_r when filled with values related by R and (2) a clause relating K_l and K_r when filled with expressions e_l and e_r for which the admissibility condition $\mathcal{T} \blacktriangleleft e_l \lesssim e_r \langle R \rangle$ holds. This condition asserts that the refinement between e_l and e_r with postcondition R is admissible under \mathcal{T} . Because R is a relation on values while return conditions in \mathcal{T} are relations on expressions, admissibility existentially quantifies over a return condition Q such that $\mathcal{T}(e_l, e_r, Q)$ holds. To connect Q with R , the definition also claims that the refinement $e'_l \lesssim e'_r \langle \mathcal{T} \rangle \{R\}$ holds for every pair of expressions e'_l and e'_r related by Q . This occurrence of the refinement relation makes its definition recursive. This explains the use of the *later modality* \triangleright [3, 41], which is one of Iris's mechanisms to introduce recursive definitions: as long as the recursive occurrences are guarded by the later modality, the definition can be constructed in Iris. The use of the persistently modality is again related to the compliance with multi-shot continuations.⁸

4.1.4 Reasoning Rules. The refinement relation enjoys a collection of powerful and high-level reasoning rules shown in Figure 3. Rules **VALUE**, **STEP-L**, and **STEP-R** are standard: Rules **STEP-L** and **STEP-R** provide the ability to partially execute code using pure reductions and Rule **VALUE** allows the user to end a refinement proof when both sides terminate with values that satisfy the postcondition. The remaining rules are novel.

⁷There is no need to enclose e_l under a universally quantified context, because *wp* already enjoys context-local reasoning.

⁸In §4.3, we show how to extend the logic with support for one-shot continuations with no changes to the model.

Rule **INTRODUCTION** states the admissibility of a refinement between e_l and e_r with postcondition R under the theory \mathcal{T} implies $e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\}$. This rule is usually applied to reason about effectful operations independently of their handlers.

Rule **EXHAUSTION** incorporates a *case-analysis* principle into the logic whereby, if $e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\}$ holds, then the derivation of $K_l[e_l] \lesssim K_r[e_r] \langle \mathcal{F} \rangle \{S\}$ splits into two subgoals: one where e_l and e_r are replaced with values related by R and another one where e_l and e_r are replaced with expressions e'_l and e'_r such that $\mathcal{T} \blacktriangleleft e'_l \lesssim e'_r \{R\}$ holds. This rule allows one to reason about handlers independently of their handlees. It is typically applied when the contexts K_l and K_r contain handlers monitoring the handlees e_l and e_r . However, it is important to note that the rule is applicable to any contexts K_l and K_r . This flexibility allows the relation of programs where handlers on both sides of the relation do not necessarily match. In §4.1.5, we return to the example of *countdown* (§2) to show this principle in action.

Rule **MONOTONICITY** enables one to weaken the postcondition R and the parameterised theory \mathcal{T} of a refinement $e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\}$. Such a reasoning principle is useful, for example, when the refinement $e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\}$ is assumed, but \mathcal{T} and R do not exactly match \mathcal{F} and S . The weakening of R to S is captured by the condition $\square \forall v_l, v_r. R(v_l, v_r) \rightarrow S(v_l, v_r)$, where the persistently modality ensures this ordering does not rely on ephemeral resources and can thus be used multiple times in case the program is reified as a multi-shot continuation and resumed multiple times. The weakening of \mathcal{T} to \mathcal{F} is captured by the *theory ordering* $\mathcal{T} \sqsubseteq \mathcal{F}$, which is similarly defined: $\square \forall e_l, e_r, Q. \mathcal{T}(e_l, e_r, Q) \rightarrow \mathcal{F}(e_l, e_r, Q)$.

Finally, Rule **BIND** enables context-local reasoning about e_l and e_r independently of their enclosing evaluation contexts K_l and K_r . The only side-condition is that there must be a theory \mathcal{T} contained in \mathcal{F} (that is, $\mathcal{T} \sqsubseteq \mathcal{F}$) that traverses the pair of contexts K_l and K_r . Informally, this says that whenever e_l and e_r are related by the theory \mathcal{T} , so are $K_l[e_l]$ and $K_r[e_r]$. The formal definition is:

$$\begin{aligned} \text{traversable}(K_l, K_r, \mathcal{T}) &\triangleq \square \forall e_l, e_r, Q. \mathcal{T}(e_l, e_r, Q) \rightarrow \\ &\exists P. \mathcal{T}(K_l[e_l], K_r[e_r], P) * \square \forall e'_l, e'_r. P(K_l[e'_l], K_r[e'_r]) \rightarrow Q(e'_l, e'_r) \end{aligned}$$

This definition is transparent to the user. In other words, when applying Rule **BIND**, the user must find a theory \mathcal{T} and prove this traversable condition. Fortunately, the predicate *traversable* works nicely in combination with the context-closure of a theory:

$$\forall \mathcal{T}, ls_l, ls_r, K_l, K_r. \text{neutral}(ls_l, K_l) \rightarrow \text{neutral}(ls_r, K_r) \rightarrow \text{traversable}(K_l, K_r, (ls_l, ls_r) \Downarrow \mathcal{T}) \quad (10)$$

Using this theorem, it is possible to derive the following version of the bind rule, where the traversable condition is replaced with more explicit conditions on the contexts K_l and K_r :

$$\text{DERIVED-BIND} \quad \frac{\text{neutral}(ls_l, K_l) \quad \text{neutral}(ls_r, K_r) \quad (ls_l, ls_r) \Downarrow \mathcal{T} \sqsubseteq \mathcal{F}}{e_l \lesssim e_r \langle (ls_l, ls_r) \Downarrow \mathcal{T} \rangle \{v_l v_r. K_l[v_l] \lesssim K_r[v_r] \langle \mathcal{F} \rangle \{R\}\} \vdash K_l[e_l] \lesssim K_r[e_r] \langle \mathcal{F} \rangle \{R\}}$$

4.1.5 Example. We briefly discuss how these rules can be used to derive Refinements 6 and 7 from §2.2. First, let us formally define the theory $\text{Timer}_{\text{spec}}^{\ell}$:

$$\begin{aligned} \text{Timer}_{\text{spec}}^{\ell} &\triangleq ([\$\text{Timer}], []) \Downarrow (\text{Get} \oplus \text{Set}) \\ \text{Get}(\text{perform } \$\text{Timer} (\text{inl } ()), !\ell, Q) &\triangleq \exists x. \ell \xrightarrow{s}^{1/2} x * \square (\ell \xrightarrow{s}^{1/2} x \rightarrow Q(x, x)) \\ \text{Set}(\text{perform } \$\text{Timer} (\text{inr } y), \ell \leftarrow y, Q) &\triangleq \ell \xrightarrow{s}^{1/2} _* * \square (\ell \xrightarrow{s}^{1/2} y \rightarrow Q((), ())) \end{aligned}$$

The definition uses the *sum operator* \oplus , which combines relations from two theories:

$$(\mathcal{T} \oplus \mathcal{F})(e_l, e_r, Q) \triangleq \mathcal{T}(e_l, e_r, Q) \vee \mathcal{F}(e_l, e_r, Q)$$

The theory *Get* allows the handlee to establish a refinement between **perform \$Timer (inl ())** and **!ℓ** in exchange for the fractional ownership $\ell \xrightarrow{s}^{1/2} x$. This assertion appears as a premise to the

return condition, which holds of the pair (x, x) . From the perspective of the handlee, this means that the fractional ownership $\ell \xrightarrow{1/2} s x$ can be reclaimed and that the expressions **perform** \$Timer (inl ()) and $!\ell$ both return x . The reading of *Set* is analogous.

Because $Timer_{spec}^\ell$ is defined using the context-closure operator, the theory can be used in conjunction with **DERIVED-BIND** (where \mathcal{F} is instantiated with $Timer_{spec}^\ell$) to carry out Refinement 6.

The key rule to establish Refinement 7 is Rule **EXHAUSTION**. It is applied using K_l instantiated with *run_st_passing*'s handler and K_r instantiated with $[]$. Furthermore, the theory \mathcal{T} in the statement of Rule **EXHAUSTION** is taken to be $Timer_{spec}^\ell$ and \mathcal{F} is instantiated with \perp . During the proof of the clause $\forall e'_l, e'_r. Timer_{spec}^\ell \blacktriangleleft e'_l \lesssim e'_r \{= \} \rightarrow K_l[e'_l] \lesssim e'_r \{= \}$, the admissibility condition gives the fractional ownership $\ell \xrightarrow{1/2} s x$ to the handler. In combination with the other assertion $\ell \xrightarrow{1/2} s x$ initially given to the handler as a premise in 7, full ownership of ℓ is claimed by the handler, which can then update ℓ in case of a set request. The return condition can be interpreted in this proof as the condition under which the handler can resume the continuation: in the case of a get request, for example, both the value x and the fractional ownership $\ell \xrightarrow{1/2} s x$ must be supplied.

4.2 blaze: A Logic for Effect Handlers with Dynamic Labels

So far, we have exclusively considered examples where the effect labels have already been allocated. This observation incites the question: how to reason about programs like *run_ask* (§3.1) where effect labels are allocated locally to avoid collision of effect names? For instance, given two clients of the ask effect *main*₁ and *main*₂ and an integer x , is it possible to establish a refinement between *run_ask* x *main*₁ and *main*₂ (**fun**_. x) when *main*₁ and *main*₂ perform arbitrary effects?

The verification of programs with local allocation of effects, such as *run_ask*, depends heavily on assumptions about fresh labels being distinct from previously allocated ones. The *baze* logic places the burden of keeping track of these assumptions entirely on the user. So, while *baze* *can* be used to reason about *run_ask*, it is not placed at the right level of abstraction. To address this limitation, we introduce *blaze*, a logic built on top of *baze* to facilitate reasoning about dynamic labels. In *blaze*, it is possible to establish (in a relatively straightforward way) a strong result about *run_ask* where assumptions about labels being distinct are hidden:

$$\forall main_1, \left(\begin{array}{l} \forall ask_1, ask_2, \mathcal{M}. \\ main_2, \left(\begin{array}{l} (\square ask_1() \lesssim_\star ask_2() \langle \mathcal{M} \rangle \{v_l v_r. v_l = v_r = x\}) \rightarrow \\ x, \mathcal{L}, R. main_1 ask_1 \lesssim_\star main_2 ask_2 \langle \mathcal{L} + \mathcal{M} \rangle \{R\} \end{array} \right) \end{array} \right) \rightarrow^* \begin{array}{l} \text{i}\text{run_ask } x \text{ main}_1 \lesssim_\star \\ \text{main}_2 (\text{fun}__. x) \langle \mathcal{L} \rangle \{R\} \end{array} \quad (11)$$

The novelty of the refinement relation $e_l \lesssim_\star e_r \langle \mathcal{L} \rangle \{R\}$ is the parameterised *list of theories* \mathcal{L} . The elements of \mathcal{L} are triples of the form $(ls_l, ls_r, \mathcal{T})$, where ls_l and ls_r are lists of effect labels respectively allocated by the implementation and the specification sides, and \mathcal{T} is a theory relating expressions that use these effects. Roughly speaking, the list \mathcal{M} in Refinement 11 is used to relate **fun**_. **perform** Ask () to **fun**_. x . Its universal quantification reflects the fact that Ask is allocated locally by *run_ask*. The list \mathcal{L} represents an ambient set of relational theories used to relate *main*₁ to *main*₂. The lists of theories \mathcal{L} and \mathcal{M} are *disjoint*, because the labels in \mathcal{L} are allocated before Ask. This assumption however exists only as part of the model of the logic. Because the model is opaque, this requirement is never directly exposed to the user.

4.2.1 Model. The formal definition of $e_l \lesssim_\star e_r \langle \mathcal{L} \rangle \{R\}$ appears in Figure 4. It unfolds to a refinement in *baze* with parameterised theory *interp*(\mathcal{L}) and premise *valid*(\mathcal{L}).

The theory *interp*(\mathcal{L}) is constructed as the iterated sum of theories $(ls_l, ls_r) \Downarrow \mathcal{T}$ for every triple $(ls_l, ls_r, \mathcal{T})$ in \mathcal{L} . In essence, this construction sacrifices the expressivity of general theories in *baze*'s refinement relation to endow the *blaze* layer with context-local reasoning by default.

$$\begin{aligned}
e_l \lesssim_{\star} e_r \langle \mathcal{L} \rangle \{R\} &\triangleq \text{valid}(\mathcal{L}) \rightarrow e_l \lesssim e_r \langle \text{interp}(\mathcal{L}) \rangle \{R\} \\
\text{interp}([]) &\triangleq \perp \quad \text{valid}(\mathcal{L}) \triangleq \\
\text{interp}((ls_l, ls_r, \mathcal{T}) :: \mathcal{L}) &\triangleq * \begin{cases} \text{distinct}(\text{labels}_i(\mathcal{L})) * \forall \$E \in \text{labels}_i(\mathcal{L}). \text{label}_i^{\square}(\$E) \\ \text{distinct}(\text{labels}_s(\mathcal{L})) * \forall \$E \in \text{labels}_s(\mathcal{L}). \text{label}_s^{\square}(\$E) \end{cases} \\
(ls_l, ls_r) \Downarrow \mathcal{T} \oplus \text{interp}(\mathcal{L}) & \\
\text{labels}_i([]) &\triangleq [] \quad \text{labels}_s([]) \triangleq [] \\
\text{labels}_i((ls_l, _, _) :: \mathcal{L}) &\triangleq ls_l + \text{labels}_i(\mathcal{L}) \quad \text{labels}_s((_, ls_r, _) :: \mathcal{L}) \triangleq ls_r + \text{labels}_s(\mathcal{L})
\end{aligned}$$

Fig. 4. Model of blaze.

The premise $\text{valid}(\mathcal{L})$ is defined using the terms $\text{labels}_i(\mathcal{L})$ and $\text{labels}_s(\mathcal{L})$, which collect the labels in \mathcal{L} that belong to the implementation side and to the specification side, respectively. The assertions $\text{label}_i^{\square}(\$E_1)$ and $\text{label}_s^{\square}(\$E_2)$ claim ownership of persistent resources obtained after the allocation of the effects $\$E_1$ and $\$E_2$.⁹ Therefore, the premise $\text{valid}(\mathcal{L})$ asserts that the labels in $\text{labels}_i(\mathcal{L})$ and in $\text{labels}_s(\mathcal{L})$ have already been allocated and are pairwise distinct. In essence, this premise represents the assumption that the theories in \mathcal{L} do not interfere with one another or new theories for newly allocated effects.

4.2.2 Reasoning rules. The reasoning rules of blaze appear in Figure 5.

Rule **EFFECT-L-★** can be used in conjunction with Rule **ADD-LABEL-L-★** to add a freshly allocated label to one of the triples in \mathcal{L} . The assertion $\text{label}_i(\$E)$ works as an exchangeable token that is forged by Rule **EFFECT-L-★** and consumed by Rule **ADD-LABEL-L-★**. Rules **EFFECT-R-★** and **ADD-LABEL-R-★** enable analogous reasoning for the specification side. The order of the triples in \mathcal{L} and the order of the labels in a triple are not important. New triples can be added with Rule **NEW-THEORY-★**.

The statement of Rule **INTRODUCTION-★** is similar to Rule **INTRODUCTION**. The theory \mathcal{T} can be chosen among any of the list \mathcal{L} . Admissibility must be shown with respect to the triple $(ls_l, ls_r, \mathcal{T})$:

$$\begin{aligned}
(ls_l, ls_r, \mathcal{T}) \blacktriangleleft e_l \lesssim_{\star} e_r \langle \mathcal{L} \rangle \{R\} &\triangleq \\
\exists e'_l, e'_r, K_l, K_r, Q. \left(\begin{array}{l} e_l = K_l[e'_l] * e_r = K_r[e'_r] * \text{neutral}(ls_l, K_l) * \text{neutral}(ls_r, K_r) * \\ \mathcal{T}(e'_l, e'_r, Q) * \square \triangleright \forall e''_l, e''_r. Q(e''_l, e''_r) \rightarrow K_l[e''_l] \lesssim_{\star} K_r[e''_r] \langle \mathcal{L} \rangle \{R\} \end{array} \right)
\end{aligned}$$

The ability to relate expressions under arbitrary contexts K_l and K_r in this definition closes the theory \mathcal{T} under neutral contexts (for ls_l and ls_r). This design choice makes it possible to state Rule **BIND-★** in a similar fashion to Rule **DERIVED-BIND** with explicit side conditions on K_l and K_r . Namely, the condition $\mathcal{L}(K_l) \subseteq \text{labels}_i(\mathcal{M})$ restricts the handlers in K_l to the effect labels in $\text{labels}_i(\mathcal{M})$. The condition $\mathcal{L}(K_r) \subseteq \text{labels}_s(\mathcal{M})$ is analogous. The condition $\mathcal{L} + \mathcal{M} \sqsubseteq_{\star} \mathcal{N}$ is defined as the multiplicity-preserving inclusion of $\mathcal{L} + \mathcal{M}$ in \mathcal{N} . It guarantees the labels in \mathcal{L} are disjoint from the labels in \mathcal{M} . In combination, these conditions restrict the handlers in K_l and K_r to not capture effects with labels from \mathcal{L} .

Finally, Rule **EXHAUSTION-★** incorporates the exhaustion principle into blaze. The expressions e_l and e_r are related under a list of theories \mathcal{M} , but the user needs to choose only one of the theories \mathcal{T} in \mathcal{M} with which to perform the case-analysis reasoning; that is, the premise requiring a relation between $K_l[e'_l]$ and $K_r[e'_r]$ assumes the admissibility with respect only to the theory \mathcal{T} . Intuitively,

⁹The assertion $\text{label}_i^{\square}(\$E_1)$ is defined in terms of a more general assertion $\text{label}_i(\$E_1, dq)$, where dq is a *discardable fraction* [58]. Taking dq as the full fraction 1 gives the assertion $\text{label}_i(\$E_1)$ whereas taking dq as the discarded fraction gives the persistent assertion $\text{label}_i^{\square}(\$E_1)$. The assertion $\text{label}_s^{\square}(\$E_2)$ is analogously defined in terms of an assertion $\text{label}_s(\$E_2, dq)$. These definitions are included in the Appendix [18, Figure 15].

EFFECT-L-★ $\triangleright \forall \$E. \text{label}_i(\$E) \rightarrow K[e\{\$E/E\}] \lesssim_\star e_r \langle \mathcal{L} \rangle \{R\}$ $K[\text{let effect } E \text{ in } e] \lesssim_\star e_r \langle \mathcal{L} \rangle \{R\}$	EFFECT-R-★ $\forall \$E. \text{label}_s(\$E) \rightarrow e_l \lesssim_\star K[e\{\$E/E\}] \langle \mathcal{L} \rangle \{R\}$ $e_l \lesssim_\star K[\text{let effect } E \text{ in } e] \langle \mathcal{L} \rangle \{R\}$
ADD-LABEL-L-★ $\text{label}_i(\$E) \quad e_l \lesssim_\star e_r \langle (\$E :: ls_l, ls_r, \mathcal{T}) :: \mathcal{L} \rangle \{R\}$ $e_l \lesssim_\star e_r \langle (ls_l, ls_r, \mathcal{T}) :: \mathcal{L} \rangle \{R\}$	ADD-LABEL-R-★ $\text{label}_s(\$E) \quad e_l \lesssim_\star e_r \langle (ls_l, \$E :: ls_l, \mathcal{T}) :: \mathcal{L} \rangle \{R\}$ $e_l \lesssim_\star e_r \langle (ls_l, ls_r, \mathcal{T}) :: \mathcal{L} \rangle \{R\}$
NEW-THEORY-★ $e_l \lesssim_\star e_r \langle ([], [], \perp) :: \mathcal{L} \rangle \{R\}$ $e_l \lesssim_\star e_r \langle \mathcal{L} \rangle \{R\}$	INTRODUCTION-★ $(ls_l, ls_r, \mathcal{T}) \in \mathcal{L} \quad (ls_l, ls_r, \mathcal{T}) \blacktriangleleft e_l \lesssim_\star e_r \langle \mathcal{L} \rangle \{R\}$ $e_l \lesssim_\star e_r \langle \mathcal{L} \rangle \{R\}$
BIND-★ $\mathcal{L}(K_l) \subseteq \text{labels}_i(\mathcal{M}) \quad \mathcal{L}(K_r) \subseteq \text{labels}_s(\mathcal{M}) \quad \mathcal{L} + \mathcal{M} \sqsubseteq_\star \mathcal{N}$ $e_l \lesssim_\star e_r \langle \mathcal{L} \rangle \{v_l v_r. K_l[v_l] \lesssim_\star K_r[v_r] \langle \mathcal{N} \rangle \{R\}\}$	$K_l[e_l] \lesssim_\star K_r[e_r] \langle \mathcal{N} \rangle \{R\}$
EXHAUSTION-★ $\mathcal{L}(K_l) \subseteq ls_l \quad \mathcal{L}(K_r) \subseteq ls_r$ $e_l \lesssim_\star e_r \langle \mathcal{M} \rangle \{R\} \quad \mathcal{M} = (ls_l, ls_r, \mathcal{T}) :: \mathcal{L} \quad \mathcal{N} = (ls_l, ls_r, \mathcal{F}) :: \mathcal{L}$ $\wedge \begin{cases} \square \forall v_l, v_r. R(v_l, v_r) \rightarrow K_l[v_l] \lesssim_\star K_r[v_r] \langle \mathcal{N} \rangle \{S\} \\ \square \forall e'_l, e'_r. (ls_l, ls_r, \mathcal{T}) \blacktriangleleft e'_l \lesssim_\star e'_r \langle \mathcal{M} \rangle \{R\} \rightarrow K_l[e'_l] \lesssim_\star K_r[e'_r] \langle \mathcal{N} \rangle \{S\} \end{cases}$	$K_l[e_l] \lesssim_\star K_r[e_r] \langle \mathcal{N} \rangle \{S\}$

Fig. 5. Reasoning rules of blaze.

this is possible because of the conditions $\mathcal{L}(K_l) \subseteq ls_l$ and $\mathcal{L}(K_r) \subseteq ls_r$, which guarantee that the remaining theories in \mathcal{M} are irrelevant in the context of K_l and K_r . The persistently modality is needed because these remaining theories can still relate effects that cause K_l or K_r to be captured in a multi-shot continuation.

4.2.3 Example. We now show how these reasoning rules can be applied to derive Refinement 11 (§4.2). The proof starts with the application of Rules **NEW-THEORY-★**, **EFFECT-L-★**, and **ADD-LABEL-L-★**, in this order. This sequence of rules has the effect of adding the fresh label **\$Ask** to a new entry in the ambient list of theories \mathcal{L} . Initially this new entry has the form $([\text{$Ask}], [], \perp)$. The core of the proof is the application of Rule **EXHAUSTION-★**, where e_l is instantiated with $\text{main}_1(\text{perform } \$Ask())$, the expression e_r with $\text{main}_2(\text{fun } _. x)$, and the theory \mathcal{T} with $\text{AskT}(\text{perform } \$Ask(), x, Q) \triangleq \square Q(x, x)$. The refinement between main_1 and main_2 directly follows from the premise of Refinement 11 with the abstract theory list \mathcal{M} instantiated with $(([\text{$Ask}], [], \text{AskT}))$. The other conditions of Rule **EXHAUSTION-★** are straightforward.

4.3 Support for One-Shot Continuations

The introduction of one-shot continuations is motivated by the fact that, in languages like OCaml, the violation of a one-shot discipline causes a runtime error. We thus follow de Vilhena and Pottier [15] and van Rooij and Krebbers [57] to introduce one-shot continuations in a way that enables the logic to rule out such runtime errors. The idea is to represent one-shot continuations with a

construct `cont` ℓK^{10} that, in addition to the reified context K , carries a location ℓ which triggers a runtime error if the continuation is resumed twice. By the soundness theorem of the logic (§4.4), a verified program either diverges or terminates, so it cannot resume a one-shot continuation twice as runtime errors are guaranteed to be absent.

Another motivation is that, as we have seen, from a logical perspective, allowing continuations to be resumed multiple times results in the addition of persistently modalities in some of the reasoning rules, most notably, in Rule **MONOTONICITY**. To provide the logic with a strong monotonicity reasoning principle applicable to fragments of code that abide by a one-shot discipline, we take inspiration from van Rooij and Krebbers [57], by introducing the *one-shot* operator $\circlearrowleft \mathcal{T}$, a semantic version of the *flip-bang* operator:¹¹

$$(\circlearrowleft \mathcal{T})(e_l, e_r, Q) \triangleq \exists P. \mathcal{T}(e_l, e_r, P) * \triangleright \forall e'_l, e'_r. P(e'_l, e'_r) \multimap Q(e'_l, e'_r)$$

The key property of this definition is that it closes a theory \mathcal{T} under a monotonicity principle on return conditions: if $(\circlearrowleft \mathcal{T})(e_l, e_r, Q)$ and $\forall v_l, v_r. Q(v_l, v_r) \multimap P(v_l, v_r)$ hold, then $(\circlearrowleft \mathcal{T})(e_l, e_r, P)$ holds. Using the notations $\circlearrowleft_{\text{ms}} \mathcal{T} \triangleq \mathcal{T}$, $\circlearrowleft_{\text{os}} \mathcal{T} \triangleq \circlearrowleft \mathcal{T}$, $\square_{\text{ms}} A \triangleq \square A$, and $\square_{\text{os}} A \triangleq A$, it is then possible to incorporate a generalised monotonicity principle into the logic:

$$\text{GEN-MONOTONICITY} \quad (\square_m \forall v_l, v_r. R(v_l, v_r) \multimap S(v_l, v_r)) \quad \begin{array}{c} e_l \lesssim e_r \langle \circlearrowleft_m \mathcal{F} \rangle \{S\} \\ \text{---} \\ e_l \lesssim e_r \langle \mathcal{T} \rangle \{R\} \quad \mathcal{T} \sqsubseteq \mathcal{F} \end{array} \vdash e_l \lesssim e_r \langle \circlearrowleft_m \mathcal{F} \rangle \{S\}$$

Taking $m = \text{ms}$ yields Rule **MONOTONICITY**, while taking $m = \text{os}$ eliminates the persistently modality, thereby allowing ephemeral resources to be used in the proof that R implies S .

The blaze logic admits a similar generalised monotonicity rule and a generalised exhaustion rule that makes use of the one-shot operator to eliminate the persistently modalities in **EXHAUSTION-★**. Both rules are included in the Appendix [18, Figure 11].

4.4 Soundness

Soundness of baze is shown by a standard *adequacy statement* [23, Thm. 7.1] that relates the notion of refinement to the underlying operational semantics of λ -blaze:¹²

THEOREM 4.1. *If $\vdash e_l \lesssim e_r \langle \perp \rangle \{\text{True}\}$, then either e_l diverges or both e_l and e_r terminate.*

Soundness of blaze follows as a corollary by $e_l \lesssim_\star e_r \langle [] \rangle \{\text{True}\} \vdash e_l \lesssim e_r \langle \perp \rangle \{\text{True}\}$.

Another common corollary of adequacy is *contextual refinement* [23, Lem. 7.2]. We cannot write the statement of contextual refinement, because it depends on types but λ -blaze is untyped. Extending λ -blaze with types is one of our directions for future work (§7).

5 Case Studies

To assess the usability of our logic, we verify refinement statements for a number of interesting effects including *concurrency* (§5.1), *Haskell-like non-determinism* (§5.2), and *state*, where, like Biernicki et al. [5, §4.2], we show state can be implemented in terms of two independent *reader* and *writer* effects. In the interest of space, we do not discuss this state effect in detail. Its implementation can be found in the Appendix [18, §C.1]. Mechanised proofs of all case studies are included in our Rocq formalisation [17].

¹⁰In our Rocq formalisation [17], we have the extended syntax $k \langle \text{as multi} \rangle^?$ for the continuation binder in an effect handler. The keywords `as multi` are optional. Their presence indicates a multi-shot semantics. Their absence indicates the handler captures the handlee in a one-shot continuation. The construct `cont` ℓK is introduced at runtime by such handlers.

¹¹This definition can also be seen as a generalisation of the *upward closure* [14, §2.2.2] to a binary setting.

¹²Theorem 4.1 states a slightly weaker result than the adequacy theorem proven in the Rocq development [17]. This weaker version, which is sufficient to show soundness, uses a fixed postcondition `True`, whereas the formalised one, which can be found in the Appendix [18, Theorem B.3], uses an arbitrary pure postcondition.

$$\begin{aligned}
 runForkSpec &\triangleq \square \forall main_1, main_2. \\
 &\left(\begin{array}{l} \forall fork_1, fork_2, \mathcal{L}. \\ forkSpec(fork_1, fork_2, \mathcal{L}) \rightarrow* \\ \{main_1 fork_1 \lesssim_\star main_2 fork_2 \langle \mathcal{L} \rangle\} \{True\} \end{array} \right) \rightarrow* \begin{array}{l} run_fork main_1 \lesssim_\star \\ main_2 (\mathbf{fun} task'. \mathbf{fork} (task'())) \{True\} \end{array} \\
 forkSpec(fork_1, fork_2, \mathcal{L}) &\triangleq \square \forall task_1, task_2. \\
 &task_1() \lesssim_\star task_2() \langle \mathcal{L} \rangle \{True\} \rightarrow* fork_1 task_1 \lesssim_\star fork_2 task_2 \langle \mathcal{L} \rangle \{True\}
 \end{aligned}$$

Fig. 6. Fork case study: Specification.

5.1 Concurrency

Effect handlers enable the implementation of *cooperative-concurrency* libraries. In such libraries, multiple tasks can be spawned and their execution is monitored by a scheduler making sure at most one task runs at a time. It is an important and interesting application of effect handlers, serving as the “*primary motivation*” for the addition of effect handlers to OCaml [38, Chapter 12, §24.5]. Here is the handler-based implementation of a *fork effect* [7, Fig. 11] in λ -blaze:

```

run_fork ≡ fun main.
  let effect Fork in let q = new_queue() in
  let run = rec run task. handle task() with
    | effect Fork task', k ⇒ push q k; run task'
    | _ ⇒ if empty q then () else (let k = pop q in k)
  in run (fun _ . main (fun task'. perform Fork task'))

```

The function *run_fork* supplies a piece of client code *main* with the functionality to fork tasks by monitoring the execution of *main* with a handler for the *Fork* effect. The handling of a *Fork* effect with payload *task'* pushes the paused continuation *k* to a queue *q*. This queue is allocated at the beginning of *run_fork*'s execution. It is initially empty, and, as an invariant, it stores continuations that can be readily resumed with *()*. Updates to *q* maintain this invariant, because, thanks to a deep-handler semantics, the continuation *k* includes the *Fork* handler at its top-most frame. After this update, the handling of *Fork* terminates by running *task'* under a new *Fork* handler. When a task terminates, if the handler finds *q* non-empty, it pops a continuation *k* from *q* representing a previously paused task and resumes the execution of this task. If *q* is empty then all scheduled tasks have executed, so the function *run_fork* terminates.

The implementation of *run_fork* is concise, but relies on advanced programming features, notably, the ability to reify contexts as first-class continuations using handlers and the ability to place these continuations in the store. The complexity of *run_fork*'s operational behaviour motivates the question: is it possible to show that the fork functionality implemented by *run_fork* can be abstracted as a real concurrent fork instruction?

In this case study, we answer this question positively by verifying in blaze that the functionality implemented by *run_fork* refines the primitive *fork* construct of λ -blaze. The formal statement is written in Figure 6. The specification of *run_fork*, the assertion *runForkSpec*, states a refinement between the application of *run_fork* to a client *main*₁ and the application of a client *main*₂ to a function **fun** *task'*. **fork** (*task'()*) that directly forks *task'*. The clients *main*₁ and *main*₂ are universally quantified in this specification. It is assumed that *main*₁ and *main*₂ can be related when respectively supplied with abstract fork implementations *fork*₁ and *fork*₂. It is the obligation of the user of the library to show the relation between *main*₁ *fork*₁ and *main*₂ *fork*₂. To establish this relation, the user can rely on a relational specification of *fork*₁ and *fork*₂, the assertion *forkSpec*,

stating a relation between the application of fork_1 to a task task_1 and the application of fork_2 to a task task_2 . To use this specification, it is again an obligation of the user to establish the relation between $\text{task}_1()$ and $\text{task}_2()$. In establishing this relation, the user can still rely on forkSpec to relate further calls to fork_1 and fork_2 in the tasks task_1 and task_2 . The refinement between main_1 and main_2 is carried out under an abstract theory list \mathcal{L} . Intuitively, this list represents the internal relational theory that is used by run_fork to relate Fork to fork . Apart from \mathcal{L} , which is abstract to the user, the specification runForkSpec assumes an *empty* ambient theory to relate the effects of main_1 and main_2 as well as the effects of two forked tasks task_1 and task_2 . In other words, the specification disallows main_1 and main_2 as well as forked tasks to perform unhandled effects. This limitation is necessary because forked tasks on the specification side of the refinement run on new empty contexts, where performing an unhandled effect constitutes a runtime error.

5.1.1 Relational Reasoning About Concurrency. Before presenting the proof of runForkSpec , we explain how we extend blaze with support for reasoning about native concurrency.¹³ The logic has support for invariants in the same way as ReLoC [23]: there are two general rules for allocating and closing invariants and one invariant-opening rule per atomic instruction. In the interest of space, we do not discuss these rules, because they are not needed in our case studies.¹⁴ Invariants are not needed, because, in all case studies, native concurrency occurs only on the specification side of the refinement, whereby, thanks to an *angelic* flavour of non-determinism, the user is (or should be) capable of deciding how threads interleave to avoid interference. Despite the substantial literature on relational concurrent separation logic [22, 23, 58, 59], we found that rules to achieve such a desirable reasoning ability are lacking with respect to three key limitations which we explain next. To address these limitations, we design *novel* relational reasoning rules for concurrency.

Limitation to refinements where forks match. In previous work (for example [23, §4.1]), it is assumed that **fork** instructions on both sides of a refinement match. This is clearly not the case for the refinement runForkSpec because only the specification side forks threads directly. To overcome this limitation, we follow Vindum et al. [59] in exposing the ghost thread-pool assertion $i \Rightarrow e$ ¹⁵ in the logic. Recall that its reading simply states thread i at the specification side runs e . Using this resource, we can split a traditional relational **fork rule** into Rules **FORK-L-★** and **FORK-R-★**, shown in Figure 7. Rule **FORK-R-★** forges a new resource $i \Rightarrow e_r$. There are many ways to spend this resource. Rule **FORK-L-★** consumes it to allow reasoning about a **fork** e_l instruction on the implementation side. As a condition to this rule, the expressions e_l and e_r must be related under the theory list \mathcal{L}^\perp , which sets every theory in \mathcal{L} to \perp . This condition guarantees the forked threads do not perform unhandled effects.

Explicit operational reasoning about thread-pool assertions. The reasoning rules introduced by Vindum et al. [59, Fig. 8] require the user to explicitly manipulate thread-pool resources; that is, the user must inspect the shape of the expression e_r in an assertion $i \Rightarrow e_r$ and select one of their rules allowing e_r to be partially executed. This is a strong limitation for the verification of runForkSpec , because the only assumption on forked tasks task_1 and task_2 is that $\text{task}_1()$ refines $\text{task}_2()$. The specific shape of task_2 is unknown. To overcome this limitation, we introduce Rule **LOGICAL-FORK-★** (Figure 7). This rule consumes a thread-pool resource $i \Rightarrow K_r[e_r]$ and, as a condition, the user must supply a subexpression e_l that refines e_r . In return, the user can reclaim the assertion $i \Rightarrow K_r[v_r]$ where e_r is replaced with its result v_r , obtained with no explicit manipulation

¹³We focus on blaze but similar reasoning principles can be achieved in baze [18, Figure 17].

¹⁴The rules for allocating, opening, and closing invariants can be found in the Appendix [18, §C.4].

¹⁵Vindum et al. [59] in fact present this resource as a *right refinement*. In our logic, the user does not explicitly manipulate this resource; it is already abstract as is, so we can keep its standard notation.

$\begin{array}{c} \text{FORK-L-★} \\ i \Rightarrow e_r \quad e_l \lesssim_\star e_r \langle \mathcal{L}^\perp \rangle \{\text{True}\} \quad K_l[\langle \rangle] \lesssim_\star e'_r \langle \mathcal{L} \rangle \{R\} \\ \hline K_l[\mathbf{fork} e_l] \lesssim_\star e'_r \langle \mathcal{L} \rangle \{R\} \end{array}$	$\begin{array}{c} \text{FORK-R-★} \\ \forall i. i \Rightarrow e_r \rightarrow e_l \lesssim_\star K_l[\langle \rangle] \langle \mathcal{L} \rangle \{R\} \\ \hline e_l \lesssim_\star K_l[\mathbf{fork} e_r] \langle \mathcal{L} \rangle \{R\} \end{array}$
$\begin{array}{c} \text{LOGICAL-FORK-★} \\ i \Rightarrow K_r[e_r] \quad e_l \lesssim_\star e_r \langle \mathcal{L}^\perp \rangle \{R\} \quad \forall v_l, v_r. R(v_l, v_r) \rightarrow i \Rightarrow K_r[v_r] \rightarrow K_l[v_l] \lesssim_\star e'_r \langle \mathcal{L} \rangle \{S\} \\ \hline K_l[e_l] \lesssim_\star e'_r \langle \mathcal{L} \rangle \{S\} \end{array}$	
$\begin{array}{c} \text{THREAD-SWAP-★} \\ i \Rightarrow K[e_r] \quad \forall j, K'. j \Rightarrow K'[e'_r] \rightarrow e_l \lesssim_\star e_r \langle \mathcal{L}^\perp \rangle \{v_l _ \exists v'_r. j \Rightarrow K'[v'_r] * R(v_l, v'_r)\} \\ \hline e_l \lesssim_\star e'_r \langle \mathcal{L} \rangle \{R\} \end{array}$	

Fig. 7. Reasoning rules for concurrency.

of the thread-pool assertion. This rule can be used in conjunction with Rule **FORK-R-★** to derive the refinement $e_1 \lesssim e'_1 \{\text{True}\} \rightarrow e_2 \lesssim e'_2 \{\text{True}\} \rightarrow e_1; e_2 \lesssim \mathbf{fork}(e'_1); e'_2 \{\text{True}\}$, which cannot be shown using the rules in Vindum et al. [59, Fig. 8] without breaking the abstraction of their refinement relation.

Access to thread-pool resource describing the main thread. With the rules discussed so far, the only way to obtain new thread-pool resources is by means of Rule **FORK-R-★**. In other words, thread-pool resources can only describe forked threads but not the *main thread* e_r on the specification side of the refinement. As we are going to see, the proof of *runForkSpec* needs access to the thread-pool resource describing the main thread. Rule **REL-SPLIT** from Vindum et al. [59, Fig. 8] supports this very feature. However, the statement relies on the fact that ReLoC’s notion of refinement $\Delta \models e_l \lesssim e_r : \tau$ is defined using $i \Rightarrow e_r$ as a premise. This makes the adaption of **REL-SPLIT** to blaze particularly difficult, because blaze’s model hides thread-pool assertions under multiple layers of abstraction.¹⁶ Instead, we introduce Rule **THREAD-SWAP-★** (Figure 7), which allows the user to trade a thread-pool resource $i \Rightarrow K[e_r]$ in exchange for a thread-pool resource $j \Rightarrow K'[e'_r]$ describing the main thread e'_r under an abstract context K' . The expression e_r becomes the new main thread on the specification side and the postcondition is updated to require the termination of e'_r , which is part of the implicit requirements of the original refinement.

5.1.2 Verification. After the allocation of an effect label **\$Fork** by *run_fork*, the crux of the proof is (1) the introduction of a relational theory *Fork* to relate **\$Fork** effects to **fork** and (2) the definition of the queue invariant in blaze. These definitions appear in Figure 8.

The theory *Fork* requires *task*₁ to refine *task*₂ as naturally expected. To allow **\$Fork** effects in *task*₁, the refinement between *task*₁ and *task*₂ assumes the theory *Fork* itself. The later modality \triangleright guards this recursive occurrence of *Fork* to facilitate the definition in Iris. The return condition asserts that both the **\$Fork** effect and **fork** return $\langle \rangle$.

Recall that, according to the informal explanation of *run_fork*, the queue stores continuations that can be readily resumed. This description is formally captured by the predicate *queueInv*. The term q represents the queue identifier and ks represents the contents of q . Concretely, ks is a list of triples $(k, (j, K))$, where k is one of the continuations in q . This connection is captured by *isQueue*($q, ks.1$), which asserts q contains the collection of continuations in ks . Because k is created by a running task that performs an effect, there must be a corresponding task on the

¹⁶The same holds for baze.

Relational theory.

$$\begin{aligned} \text{Fork}(\text{perform } \$\text{Fork } \text{task}_1, \text{ fork } (\text{task}_2()), Q) &\triangleq \\ &\triangleright \text{task}_1() \lesssim_\star \text{task}_2() \langle \{([\$\text{Fork}], [], \text{Fork})\} \{ \text{True} \} * Q((), ()) \rangle \end{aligned}$$

Invariants and predicates.

$$\begin{aligned} \text{queueInv}(q, ks, ks') &\triangleq \text{isQueue}(q, ks.1) * \\ &(*_{(k, (j, K)) \in ks} \exists e_r. j \Rightarrow K[e_r] * \text{ready}(q, k(), e_r)) * (*_{(_, (j, K)) \in ks'} \exists v_r. j \Rightarrow K[v_r]) \\ \text{ready}(q, e_l, e_r) &\triangleq \forall ks, ks'. \triangleright \text{queueInv}(q, ks, ks') -* \\ e_l \lesssim_\star e_r \langle \{([\$\text{Fork}], [], \perp)\} \{ \text{queueInv}(q, [], ks \#+ ks') \} \rangle \end{aligned}$$

Fig. 8. Fork case study: Internal logical definitions.

specification side that k refines. The thread identifier j and the context K are used to describe the state of this task: it is an expression e_r such that $j \Rightarrow K[e_r]$. Finally, the term ks' in queueInv represents the terminated tasks on the specification side once used to describe continuations in ks .

When handling a $\$Fork$ effect with payload task'_1 , the specification side is a program of the form $K_r[\text{fork } (\text{task}'_2())]$. After the application of Rule **FORK-R-★**, the newly obtained resource $i \Rightarrow \text{task}'_2()$ is traded, via Rule **THREAD-SWAP-★**, for a thread-pool resource $j \Rightarrow K'[K_r(())]$ describing the main thread. This resource is used to show the queue invariant is preserved after pushing k . The proof then carries on with $\text{run task}'_1$ on the implementation side and the specification side correctly adjusted to $\text{task}'_2()$. Upon termination of a task, if the queue is non-empty, a continuation k is taken from the queue. At this point, Rule **LOGICAL-FORK-★** is used in conjunction with the thread-pool resource and the ready assumption retrieved from the queue invariant, thus concluding the proof.

5.1.3 Async/Await. We prove a similar refinement for an *asynchronous-computation* library offering *async* and *await* effects [15, 19]. The implementation run_coop_1 (Figure 9) is the translation to λ -blaze of the OCaml implementation from Dolan et al. [19, Fig. 1].

In addition to a queue of ready continuations, run_coop_1 also stores continuations in *promises*. Abstractly, a promise p represents the result of a running task. Continuations in p wait for this result. They can be readily resumed once the task finishes, when they are transferred to the queue. We show run_coop_1 refines run_coop_2 , which implements *async* more directly using **fork** instead of a queue of continuations. The implementation of *await* by run_coop_2 still relies on a handler and also uses promises to manage waiting threads. To avoid races on promises, run_coop_2 uses locks.

The proof that run_coop_1 refines run_coop_2 relies on a queue invariant similar to queueInv (Figure 8). Other internal definitions are adapted from de Vilhena and Pottier [15] (who verify a similar library in a unary setting in Iris). All definitions are included in the Appendix [18, §C.2.1].

Finally, we also prove the negative result that run_coop_1 does not refine the following handler-free implementation of *async* and *await* by run_coop_3 , where *async* is implemented using **fork** and *await* is implemented by *busy waiting*:

$$\begin{aligned} \text{deadlock} &\triangleq \text{fun } \text{async } \text{await}. & \text{run_coop}_3 &\triangleq \text{fun } \text{main}. \\ \text{let } r = \text{ref } (\text{inl } ()) \text{ in} & & \text{let } \text{async} = \text{fun } \text{task}. \\ \text{let } p = \text{async } (\text{rec } f () . & & \text{let } p = \text{ref } (\text{inl } ()) \text{ in} \\ \quad \text{match } !r \text{ with} & & \quad \text{fork } (\text{let } y = \text{task}() \text{ in } p \leftarrow (\text{inr } y)); p \\ \quad | \text{inl } () \Rightarrow \text{async } (\text{fun } _. ()); f () & & \quad \text{inlet } \text{await} = \text{rec } \text{await } p. \\ \quad | \text{inr } p \Rightarrow \text{await } p & & \quad \text{match } !p \text{ with } \text{inl } () \Rightarrow \text{await } p \mid \text{inr } v \Rightarrow v \\ \quad) \text{ in } r \leftarrow \text{inr } p; \text{ await } p & & \quad \text{in } \text{main } \text{async } \text{await} \end{aligned}$$

```

 $\text{run\_coop}_1 \triangleq \text{fun main.}$ 
 $\quad \text{let effect Coop in}$ 
 $\quad \text{let } q = \text{new\_queue}() \text{ in}$ 
 $\quad \text{let } next = \text{fun }_. \quad$ 
 $\quad \quad \text{if empty } q \text{ then } () \text{ else } (\text{pop } q)()$ 
 $\quad \text{in}$ 
 $\quad \text{let } run = \text{rec } run p \text{ task. handle task() with}$ 
 $\quad | \text{effect Coop request, } k \Rightarrow$ 
 $\quad \quad \text{match request with}$ 
 $\quad \quad | \text{inl } task' \Rightarrow$ 
 $\quad \quad \quad \text{let } p' = \text{ref } (\text{inr } []) \text{ in}$ 
 $\quad \quad \quad \text{push } q (\text{fun }_. \text{ k } p'); run p' \text{ task'}$ 
 $\quad \quad | \text{inr } p' \Rightarrow \text{match } !p' \text{ with}$ 
 $\quad \quad \quad | \text{inl } x \Rightarrow k x$ 
 $\quad \quad \quad | \text{inr } ks \Rightarrow p' \leftarrow \text{inr } (k :: ks); next()$ 
 $\quad | y \Rightarrow$ 
 $\quad \quad \text{let } (\text{inr } ks) = !p \text{ in } p \leftarrow \text{inl } y;$ 
 $\quad \quad \text{iter } (\text{fun } k. \text{ push } q (\text{fun }_. \text{ k } y)) \text{ ks;}$ 
 $\quad \quad next()$ 
 $\quad \text{in}$ 
 $\quad \text{let } async = \text{fun task'. perform Coop (inl task') in}$ 
 $\quad \text{let } await = \text{fun } p'. \text{ perform Coop (inr } p') \text{ in}$ 
 $\quad \text{let } p = \text{ref } (\text{inr } []) \text{ in}$ 
 $\quad run p (\text{fun }_. \text{ main } async \text{ await})$ 

 $\text{run\_coop}_2 \triangleq \text{fun main.}$ 
 $\quad \text{let effect Await in}$ 
 $\quad \text{let } new\_promise = \text{fun }_. \quad$ 
 $\quad \quad (\text{ref } (\text{inr } []), \text{new\_lock}())$ 
 $\quad \text{in}$ 
 $\quad \text{let } run = \text{rec } run p \text{ task.}$ 
 $\quad \text{handle task() with}$ 
 $\quad | \text{effect Await } p', k \Rightarrow$ 
 $\quad \quad \text{acquire } p'.2; \text{match } !p'.1 \text{ with}$ 
 $\quad \quad | \text{inl } x \Rightarrow \text{release } p'.2; k x$ 
 $\quad \quad | \text{inr } ks \Rightarrow p'.1 \leftarrow \text{inr } (k :: ks);$ 
 $\quad \quad \text{release } p'.2$ 
 $\quad | y \Rightarrow \text{acquire } p.2;$ 
 $\quad \quad \text{let } (\text{inr } ks) = !p.1 \text{ in}$ 
 $\quad \quad p.1 \leftarrow \text{inl } y; \text{release } p.2;$ 
 $\quad \quad \text{iter } (\text{fun } k. \text{ fork } (k y)) \text{ ks}$ 
 $\quad \text{in}$ 
 $\quad \text{let } async = \text{fun task'.}$ 
 $\quad \quad \text{let } p' = new\_promise() \text{ in}$ 
 $\quad \quad \text{fork } (run p' \text{ task'}); p'$ 
 $\quad \text{in}$ 
 $\quad \text{let } await = \text{fun } p'. \text{ perform Await } p' \text{ in}$ 
 $\quad \text{let } p = new\_promise() \text{ in}$ 
 $\quad run p (\text{fun }_. \text{ main } async \text{ await})$ 

```

Fig. 9. Async/await implementations.

The key idea is to adapt the *deadlock* example from de Vilhena [14, Fig. 4.2] to exhibit a client that terminates when using the handler-based library but diverges otherwise.¹⁷ In short, the client *deadlock* creates a cyclic dependency between p and itself. With the implementation of *async* and *await* by run_coop_3 , when *deadlock* executes the final instruction *await p*, it diverges, because p is never fulfilled. With the implementation of *async* and *await* by run_coop_1 , on the other hand, when *deadlock* executes the final instruction *await p*, it is captured in a continuation and stored in p . The internal queue managed by run_coop_1 becomes empty, so it terminates.

5.2 Algebraic Effects: Haskell-Like Non-Determinism

In this case study, we are interested in evaluating how relational theories can be used to reason about *algebraic effects* [44]. As an illustration, we consider the pair of constructs **or** and **fail**, where $e_1 \text{ or } e_2$ models the functionality to non-deterministically run e_1 or e_2 , and **fail** represents a failed execution path. These constructs are written in λ -blaze using a global effect $\$ND$: $e_1 \text{ or } e_2 \triangleq (\text{perform } \$ND (\text{inl } (\text{fun }_. \text{ e}_1, \text{ fun }_. \text{ e}_2)))()$ and $\text{fail} \triangleq \text{perform } \$ND (\text{inr } ())$. The construct $e_1 \text{ or } e_2$ performs a $\$ND$ effect with thunked versions of e_1 and e_2 . After one of them is non-deterministically chosen by the handler, its execution is forced with $()$. The construct **fail** just performs a $\$ND$ effect.

Plotkin and Pretnar [46] show that **or** and **fail** can be described by the *algebraic theory* of a *monoid*: $e_1 \text{ or } (e_2 \text{ or } e_3) = (e_1 \text{ or } e_2) \text{ or } e_3$ and $e \text{ or fail} = \text{fail or } e = e$. Such an algebraic theory

¹⁷The precise statement is included in the Appendix [18, §C.2.2].

```

 $\text{run\_nd\_pure} \triangleq \text{fun main.}$ 
 $\quad \text{handle main() with}$ 
 $\quad | \text{effect } \$\text{ND request, } k \Rightarrow$ 
 $\quad | \text{match request with}$ 
 $\quad | \text{inl } (t_1, t_2) \Rightarrow k \ t_1 + k \ t_2$ 
 $\quad | \text{inr } () \Rightarrow []$ 
 $\quad | y \Rightarrow [y]$ 

 $\text{run\_nd\_rand} \triangleq \text{fun main. handle main() with}$ 
 $\quad | \text{effect } \$\text{ND request, } k \Rightarrow$ 
 $\quad | \text{match request with}$ 
 $\quad | \text{inl } (t_1, t_2) \Rightarrow \text{let } b = \text{ref true in}$ 
 $\quad | \quad \text{fork } (b \leftarrow \text{false}); \text{if } !b \text{ then } k \ t_1 \text{ else } k \ t_2$ 
 $\quad | \text{inr } () \Rightarrow (\text{rec } f(). f())()$ 
 $\quad | y \Rightarrow y$ 

```

Fig. 10. Non-determinism handlers.

can be used not only to reason about **or** and **fail** but also to state the correctness of an effect handler providing an implementation of these effects. In short, a handler is correct when the handling of two programs, that are equal according to the algebraic theory, yields equal results.

This equational correctness criterion suits a pure setting well, but precludes its application to cases where the effects **or** and **fail** are implemented using native non-determinism. For example, consider the two handler implementations that appear in Figure 10. The implementation of $e_1 \text{ or } e_2$ provided by *run_nd_pure* uses a list to collect the results of returning e_1 and the results of returning e_2 . Paths signalled by **fail** are not added to this list.¹⁸ The implementation of $e_1 \text{ or } e_2$ provided by *run_nd_rand* chooses the expression to run by reading a location b that holds **true** initially but is non-deterministically set to **false** by a forked thread.¹⁹ The handling of **fail** diverges.

The correctness criterion of Plotkin and Pretnar [46] can be used to justify *run_nd_pure* provides a correct implementation of **or** and **fail** with respect to their algebraic theory. However, *run_nd_rand* falls out of the scope of their approach. Using relational theories of blaze, it is possible to introduce a similar handler-correctness criterion applicable to both *run_nd_pure* and *run_nd_rand*:

$$\text{runNdCorrect}(\text{run}) \triangleq \forall \text{main}_1, \text{main}_2. \begin{cases} \text{main}_1() \lesssim_{\star} \text{main}_2() \langle ([\$ND], [\$ND], Nd) :: \mathcal{L} \rangle \{=}\ -* \\ \text{main}_2, \mathcal{L}. \langle \text{run main}_1 \lesssim_{\star} \text{run main}_2 \langle ([\$ND], [\$ND], \perp) :: \mathcal{L} \rangle \{=} \end{cases}$$

The predicate *runNdCorrect(run)* asserts the correctness of a handler *run* with respect to a relational theory *Nd*. It states that the handling of two handlees *main*₁ and *main*₂ yields the same results assuming *main*₁ and *main*₂ are related under the theory *Nd* for **ND**. The handler *run* cannot itself rely on **ND** and it must not intercept other effects related by \mathcal{L} . The theory *Nd* enables algebraic reasoning about **or** and **fail**. It is written as the sum of several theories expressing their algebraic laws: $Nd \triangleq \text{Assoc}_1 \oplus \text{Assoc}_2 \oplus \text{Unit}_1 \oplus \text{Unit}_2 \oplus \text{Unit}_3 \oplus \text{Unit}_4 \oplus \text{Refl}_1 \oplus \text{Refl}_2$, where

$$\text{Assoc}_1(e_{11} \text{ or } (e_{12} \text{ or } e_{13}), (e_{21} \text{ or } e_{22}) \text{ or } e_{23}, Q) \triangleq \square Q(e_{11}, e_{21}) * \square Q(e_{12}, e_{22}) * \square Q(e_{13}, e_{23})$$

$$\text{Assoc}_2((e_{11} \text{ or } e_{12}) \text{ or } e_{23}, e_{21} \text{ or } (e_{22} \text{ or } e_{23}), Q) \triangleq \square Q(e_{11}, e_{21}) * \square Q(e_{12}, e_{22}) * \square Q(e_{13}, e_{23})$$

$$\text{Unit}_1(e_1 \text{ or fail}, e_2, Q) \triangleq \text{Unit}_2(\text{fail or } e_1, e_2, Q) \triangleq \square Q(e_1, e_2)$$

$$\text{Unit}_3(e_1, e_2 \text{ or fail}, Q) \triangleq \text{Unit}_4(e_1, \text{fail or } e_2, Q) \triangleq \text{f1} * \square Q(e_1, e_2)$$

$$\text{Refl}_1(e_{11} \text{ or } e_{12}, e_{21} \text{ or } e_{22}, Q) \triangleq \square Q(e_{11}, e_{21}) * \square Q(e_{12}, e_{22})$$

$$\text{Refl}_2(\text{fail, fail, } _) \triangleq \text{True}.$$

The theory *Assoc*₁ captures the associativity of **or**. The return condition *Q* is used to express the condition that the relation holds up to a relation of the subexpressions. The other theories are written in a similar style, except for *Unit*₃ and *Unit*₄, which charge the user one *later credit* [50], part of Iris's machinery to avoid cyclic proofs. Without the charge of one later credit, the theory *Unit*₄, for

¹⁸This implementation is similar to the list instance of MonadPlus's *mplus* and *mzero* [25].

¹⁹This implementation is originally given by Frumin et al. [23, §6.4].

example, could be used to relate a terminating e_1 to a diverging e_2 such as $(\mathbf{rec} f(). \mathbf{fail} \mathbf{or} f())()$. This claim is formally proved [17].

As noted in §1, blaze cannot express algebraic theories that are closed under transitivity. Therefore, Nd is symmetric and reflexive, but not transitive. The lack of support for transitivity is a known limitation of step-indexed relational logics [8, 27]. Nd is however sufficiently expressive to relate non-trivial examples of handlees [18, §C.3]. Using the $\text{run}_{Nd}\text{Correct}$ correctness criterion, we show that both run_nd_pure and run_nd_rand are correct with respect to Nd : $\text{run}_{Nd}\text{Correct}(\text{run_nd_pure})$ and $\text{run}_{Nd}\text{Correct}(\text{run_nd_rand})$ hold.

6 Related Work

To our knowledge, this is the first work to introduce a relational separation logic for effect handlers. In the following paragraphs, we discuss work within closely related topics.

Relational reasoning about effect handlers. Building on the notion of *algebraic effects*, where an effect is described by an equational theory, Plotkin and Pretnar [46] introduce the notion of correctness of handlers whereby the handler of an effect is correct if the handler *respects* the equations describing this effect. This equational approach is well-suited to strictly functional programs but has never been extended to languages with concurrency and mutable state. We follow a different approach, namely relational separation logic, but take inspiration from equational reasoning to introduce a notion of handler correctness that supports these features (§5.2).

Biernacki et al. [5] introduce binary logical relations for effect handlers. Their *biorthogonal-closed* [43] style of relations inspires similar definitions by several authors [7, 40, 60]. Such binary logical relations can be used as an intermediary step in the proof of contextual refinement. Biernacki et al. [5] explore this approach to establish interesting examples of refinement, including a statement about the *ask effect* [5, §4.1], similar to the one studied in §4.2.3, and one about the *state effect* [5, §4.2], ported to our system in our Rocq formalisation.

Logical relations can be used to develop high-level reasoning principles. Biernacki et al. [5]’s Lemma 2, for example, can be seen as a form of bind rule. The main limitation of previous logical-relations approaches is the lack of a comprehensive set of such high-level reasoning rules with which the user can verify relational properties of programs with handlers without ever being exposed to details of the model of the logic. Using separation logic as the foundation of our logic also has the advantage of having a richer assertion language than a language limited to the interpretation of syntactic types. Such expressivity is key in adding support for higher-order store and concurrency. (Even though the latter, as discussed in §5.1.1, required original work.)

Relational reasoning about continuations in Iris. Timany and Birkedal [55] devise binary logical relations for programs that manipulate *undelimited continuations* captured by `callcc`. They use these logical relations to verify multiple challenging examples of refinement, one of which is similar to the fork library we verify in §5.1. Namely, they show that the `callcc`-based implementation of fork written in a sequential language refines the `fork` construct of a language with native cooperative concurrency. The refinement therefore relates programs written in different languages. To carry out this refinement, they devise *cross-language* logical relations. Like previous works exploiting logical relations, and unlike our work, the lack of a comprehensive set of high-level reasoning rules necessitates the proofs to be carried out at the level of Iris’s weakest precondition `wp`, which, in their setting, is inconvenient because, in the presence of `callcc`, the bind rule for their version of `wp` is unsound. They mitigate this inconvenience by introducing the *context-local weakest precondition*, which admits the bind rule for the price of reduced support for `callcc`. (Although notions of weakest precondition that admit the bind rule while keeping convenient support for `callcc` exist [14, §6.3.2].)

Relational theories. de Vilhena and Pottier [15] introduce *protocols* as a mechanism to allow modular reasoning about programs with effect handlers in a unary setting. The domain of relational theories *iThy* (§4.1.1) can be seen as a generalisation to a binary setting of the domain of protocols [15, Fig. 4] ($Val \rightarrow (Val \rightarrow iProp) \rightarrow iProp$). An immediate generalisation is to replace *Val* with a binary type $Val \times Val$. A more subtle generalisation is to subsequently replace *Val* with *Expr*. This is needed to allow relations between effectful and non-effectful expressions. For the same reason, Biernacki et al. [5] introduce a similar domain of *semantic effects* *Eff* [5, §3.2], defined as a predicate of type $(Expr^2 \times (Expr^2 \rightarrow SProp)) \rightarrow SProp$, where *SProp* is a type of *step-indexed assertions*. Allain et al. [2] introduce a domain of protocols in Iris that coincides exactly with *iThy*. However, their focus is on the proof of correctness of compiler optimisations in a fragment of OCaml without handlers. Consequently, they derive a notion of simulation that admits a general bind rule with no conditions on contexts. To validate this rule, their simulation relation, by default, closes protocols under arbitrary evaluation contexts. In *baze*, we opt for a more flexible context-local reasoning principle where the user can choose when and under which contexts to close theories via the context-closure operator (§4.1.2). This flexibility is key in the layered construction of *blaze*.

Reasoning about dynamic labels. de Vilhena and Pottier [16] introduce *TesLogic*, a unary logic for effect handlers with dynamic labels in a language similar to λ -*blaze*. The model of *blaze* is inspired by how *TesLogic* builds on top of *Hazel* [15], a unary logic for handlers which, like *baze*, lacks the abstraction principles for dynamic labels. The rules of *TesLogic* [14, Fig. 7.2], however, differ from the ones in *blaze* in key ways: whereas they have an explicit rule to reason about handlers, Rule *EXHAUSTION-★* can be applied to contexts without handlers; and, whereas their bind rule is limited to neutral contexts, Rule *BIND-★* can be applied to contexts with handlers.

Flexible relational reasoning rules for concurrency. Like Vindum et al. [59], we notice limitations of the reasoning rules for concurrency provided by standard relational separation logic. We have already compared the differences between our approaches in §5.1.1. In short, we both rely on *ghost thread-pool* assertions $i \Rightarrow e$ describing the state of thread i on the specification side. However, while their rules require the user to explicitly execute e , our rules use the assertions $i \Rightarrow e$ merely as tokens that can be forged, spent, or exchanged during the construction of a proof.

7 Future Work

Limitations of our current framework indicate directions for future work. An important deficiency is the lack of a type system. In a relational setting, a type system is particularly useful, because it offers a syntax-directed approach to prove refinements of the form $e \lesssim e$. In the future, we would like to remedy this deficiency by extending λ -*blaze* with a type system for handlers with dynamic labels, such as *Tes* [16]. It would be interesting to see how *blaze* could be used to devise a binary-logical-relations interpretation of *Tes* and whether the resulting interpretation could be used to show *Tes* enforces abstraction principles for programming with handlers, such as the *absence of accidental handling* [60, 61]. Finally, we would like to explore alternative definitions of the model. We suspect the later modality in the definition of *baze*'s refinement relation can be eliminated by using an alternative method for constructing recursive definitions, namely Iris's *greatest fixpoint operator* [36, 53]. Following recent work [2, 24], we would also like to investigate the implications of generalising the type of postconditions to a predicate on pairs of expressions. We believe this generalisation could improve context-local reasoning by allowing our bind rules (*BIND* and *BIND-★*) to focus on pairs of expressions that do not necessarily terminate synchronously.

Data Availability

The Rocq formalisation is available at github.com/DeVilhena-Paulo/blaze. A snapshot of the Rocq formalisation is persistently available at doi.org/10.5281/zenodo.17308673.

Acknowledgments

We would like to thank François Pottier for putting the authors in touch, thus enabling this collaboration. Furthermore, we would like to thank the anonymous paper and artifact reviewers for their insightful suggestions. The first author is supported by the UKRI Future Leaders Fellowship MR/V024299/1. The second and fourth authors are supported, in part, by ERC grant COCONUT (grant no. 101171349), funded by the European Union; and grant *Cyclic Structures in Programs and Proofs* (file number OCENW.XL.23.089), funded by the Dutch Research Council (NWO). Views and opinions expressed are however those of the author(s) only and do not necessarily reflect those of the European Union or the European Research Council Executive Agency. Neither the European Union nor the granting authority can be held responsible for them.

References

- [1] Amal J. Ahmed, Andrew W. Appel, and Roberto Virga. 2002. A Stratified Semantics of General References Embeddable in Higher-Order Logic. In *Logic in Computer Science (LICS)*. 75–86. [doi:10.1109/LICS.2002.1029818](https://doi.org/10.1109/LICS.2002.1029818)
- [2] Clément Allain, Frédéric Bour, Basile Clément, François Pottier, and Gabriel Scherer. 2025. Tail Modulo Cons, OCaml, and Relational Separation Logic. *Proceedings of the ACM on Programming Languages* 9, POPL (2025), 2337–2363. [doi:10.1145/3704915](https://doi.org/10.1145/3704915)
- [3] Andrew W. Appel, Paul-André Melliès, Christopher D. Richards, and Jérôme Vouillon. 2007. A very modal model of a modern, major, general type system. In *Principles of Programming Languages (POPL)*. 109–122. [doi:10.1145/1190216.1190235](https://doi.org/10.1145/1190216.1190235)
- [4] Andrej Bauer and Matija Pretnar. 2014. An Effect System for Algebraic Effects and Handlers. *Logical Methods in Computer Science* 10, 4 (2014). [doi:10.2168/LMCS-10\(4:9\)2014](https://doi.org/10.2168/LMCS-10(4:9)2014)
- [5] Dariusz Biernacki, Maciej Piróg, Piotr Polesiuk, and Filip Sieczkowski. 2018. Handle with care: relational interpretation of algebraic effects and handlers. *Proceedings of the ACM on Programming Languages* 2, POPL (2018), 8:1–8:30. [doi:10.1145/3158096](https://doi.org/10.1145/3158096)
- [6] Dariusz Biernacki, Maciej Piróg, Piotr Polesiuk, and Filip Sieczkowski. 2019. Abstracting algebraic effects. *Proceedings of the ACM on Programming Languages* 3, POPL (2019), 6:1–6:28. [doi:10.1145/3290319](https://doi.org/10.1145/3290319)
- [7] Dariusz Biernacki, Maciej Piróg, Piotr Polesiuk, and Filip Sieczkowski. 2020. Binders by day, labels by night: effect instances via lexically scoped handlers. *Proceedings of the ACM on Programming Languages* 4, POPL (2020), 48:1–48:29. [doi:10.1145/3371116](https://doi.org/10.1145/3371116)
- [8] Lars Birkedal and Aleš Bizjak. 2012. A note on the transitivity of step-indexed logical relations. (Nov. 2012). <https://abizjak.github.io/documents/notes/step-indexed-transitivity.pdf>
- [9] John Boyland. 2003. Checking Interference with Fractional Permissions. In *Static Analysis Symposium (SAS) (Lecture Notes in Computer Science, Vol. 2694)*. 55–72. [doi:10.1007/3-540-44898-5_4](https://doi.org/10.1007/3-540-44898-5_4)
- [10] Jonathan Immanuel Brachthäuser, Philipp Schuster, and Klaus Ostermann. 2020. Effekt: Capability-passing style for type- and effect-safe, extensible effect handlers in Scala. *Journal of Functional Programming* 30 (2020), e8. [doi:10.1017/S0956796820000027](https://doi.org/10.1017/S0956796820000027)
- [11] Edwin C. Brady. 2013. Programming and reasoning with algebraic effects and dependent types. In *International Conference on Functional Programming (ICFP)*. 133–144. [doi:10.1145/2500365.2500581](https://doi.org/10.1145/2500365.2500581)
- [12] Ezra Cooper, Sam Lindley, Philip Wadler, and Jeremy Yallop. 2006. Links: Web Programming Without Tiers. In *Formal Methods for Components and Objects (Lecture Notes in Computer Science, Vol. 4709)*. 266–296. [doi:10.1007/978-3-540-74792-5_12](https://doi.org/10.1007/978-3-540-74792-5_12)
- [13] Ana Lúcia de Moura and Roberto Ierusalimschy. 2009. Revisiting Coroutines. *ACM Transactions on Programming Languages and Systems* 31, 2 (2009), 1–31. [doi:10.1145/1462166.1462167](https://doi.org/10.1145/1462166.1462167)
- [14] Paulo Emílio de Vilhena. 2022. *Proof of Programs with Effect Handlers*. Ph.D. Dissertation. Université Paris Cité. <https://inria.hal.science/tel-03891381>
- [15] Paulo Emilio de Vilhena and François Pottier. 2021. A Separation Logic for Effect Handlers. *Proceedings of the ACM on Programming Languages* 5, POPL (2021). [doi:10.1145/3434314](https://doi.org/10.1145/3434314)
- [16] Paulo Emílio de Vilhena and François Pottier. 2023. A Type System for Effect Handlers and Dynamic Labels. In *European Symposium on Programming (ESOP) (Lecture Notes in Computer Science, Vol. 13990)*. 225–252. [doi:10.1007/978-3-031-22800-0_9](https://doi.org/10.1007/978-3-031-22800-0_9)

3-031-30044-8_9

- [17] Paulo Emilio de Vilhena, Simcha van Collem, Ines Wright, and Robbert Krebbers. 2026. A Relational Separation Logic for Effect Handlers – Rocq Formalisation. <https://github.com/DeVilhena-Paulo/blaze>. doi:[10.5281/zenodo.17308673](https://doi.org/10.5281/zenodo.17308673)
- [18] Paulo Emilio de Vilhena, Simcha van Collem, Ines Wright, and Robbert Krebbers. 2026. A Relational Separation Logic for Effect Handlers – Technical Appendix. <https://devilhena-paulo.github.io/files/blaze.pdf>. doi:[10.5281/zenodo.17308673](https://doi.org/10.5281/zenodo.17308673)
- [19] Stephen Dolan, Spiros Eliopoulos, Daniel Hillerström, Anil Madhavapeddy, K. C. Sivaramakrishnan, and Leo White. 2017. Concurrent System Programming with Effect Handlers. In *Trends in Functional Programming (TFP) (Lecture Notes in Computer Science, Vol. 10788)*. 98–117. doi:[10.1007/978-3-319-89719-6_6](https://doi.org/10.1007/978-3-319-89719-6_6)
- [20] Derek Dreyer, Amal Ahmed, and Lars Birkedal. 2011. Logical Step-Indexed Logical Relations. *Logical Methods in Computer Science* 7, 2 (2011). doi:[10.2168/LMCS-7\(2:16\)2011](https://doi.org/10.2168/LMCS-7(2:16)2011)
- [21] Ivana Filipovic, Peter W. O’Hearn, Noam Rinetzky, and Hongseok Yang. 2010. Abstraction for concurrent objects. *Theoretical Computer Science* 411, 51–52 (2010), 4379–4398. doi:[10.1016/J.TCS.2010.09.021](https://doi.org/10.1016/J.TCS.2010.09.021)
- [22] Dan Frumin, Robbert Krebbers, and Lars Birkedal. 2018. ReLoC: A Mechanised Relational Logic for Fine-Grained Concurrency. In *Logic in Computer Science (LICS)*. 442–451. doi:[10.1145/3209108.3209174](https://doi.org/10.1145/3209108.3209174)
- [23] Dan Frumin, Robbert Krebbers, and Lars Birkedal. 2021. ReLoC Reloaded: A Mechanized Relational Logic for Fine-Grained Concurrency and Logical Atomicity. *Logical Methods in Computer Science* 17, 3 (2021). doi:[10.46298/LMCS-17\(3:9\)2021](https://doi.org/10.46298/LMCS-17(3:9)2021)
- [24] Lennard Gähler, Michael Sammller, Simon Spies, Ralf Jung, Hoang-Hai Dang, Robbert Krebbers, Jeehoon Kang, and Derek Dreyer. 2022. Simuliris: a separation logic framework for verifying concurrent program optimizations. *Proceedings of the ACM on Programming Languages* 6, POPL (2022), 1–31. doi:[10.1145/3498689](https://doi.org/10.1145/3498689)
- [25] Haskell Community. 2023. Alternative and MonadPlus. https://en.wikibooks.org/wiki/Haskell/Alternative_and_MonadPlus
- [26] Daniel Hillerström and Sam Lindley. 2016. Liberating effects with rows and handlers. In *International Workshop on Type-Driven Development (TyDe@ICFP)*. 15–27. doi:[10.1145/2976022.2976033](https://doi.org/10.1145/2976022.2976033)
- [27] Chung-Kil Hur, Derek Dreyer Georg, Neis, and Viktor Vafeiadis. 2012. The marriage of bisimulations and Kripke logical relations. In *Principles of Programming Languages (POPL)*. 59–72. doi:[10.1145/2103656.2103666](https://doi.org/10.1145/2103656.2103666)
- [28] Ralf Jung, Robbert Krebbers, Lars Birkedal, and Derek Dreyer. 2016. Higher-order ghost state. In *International Conference on Functional Programming (ICFP)*. 256–269. doi:[10.1145/2951913.2951943](https://doi.org/10.1145/2951913.2951943)
- [29] Ralf Jung, Robbert Krebbers, Jacques-Henri Jourdan, Aleš Bizjak, Lars Birkedal, and Derek Dreyer. 2018. Iris from the ground up: A modular foundation for higher-order concurrent separation logic. *Journal of Functional Programming* 28 (2018), e20. doi:[10.1017/S0956796818000151](https://doi.org/10.1017/S0956796818000151)
- [30] Ralf Jung, David Swasey, Filip Sieczkowski, Kasper Svendsen, Aaron Turon, Lars Birkedal, and Derek Dreyer. 2015. Iris: monoids and invariants as an orthogonal basis for concurrent reasoning. In *Principles of Programming Languages (POPL)*. 637–650. doi:[10.1145/2676726.2676980](https://doi.org/10.1145/2676726.2676980)
- [31] Ohad Kammar, Paul B. Levy, Sean K. Moss, and Sam Staton. 2017. A monad for full ground reference cells. In *Logic in Computer Science (LICS)*. doi:[10.1109/LICS.2017.8005109](https://doi.org/10.1109/LICS.2017.8005109)
- [32] Oleg Kiselyov and Hiromi Ishii. 2015. Freer monads, more extensible effects. In *Proceedings of the 2015 ACM SIGPLAN Symposium on Haskell (Haskell ’15)*. 94–105. doi:[10.1145/2804302.2804319](https://doi.org/10.1145/2804302.2804319)
- [33] Robbert Krebbers, Jacques-Henri Jourdan, Ralf Jung, Joseph Tassarotti, Jan-Oliver Kaiser, Amin Timany, Arthur Charguéraud, and Derek Dreyer. 2018. MoSeL: a general, extensible modal framework for interactive proofs in separation logic. *Proceedings of the ACM on Programming Languages* 2, ICFP (2018), 77:1–77:30. doi:[10.1145/3236772](https://doi.org/10.1145/3236772)
- [34] Robbert Krebbers, Ralf Jung, Aleš Bizjak, Jacques-Henri Jourdan, Derek Dreyer, and Lars Birkedal. 2017. The essence of higher-order concurrent separation logic. In *European Symposium on Programming (ESOP) (Lecture Notes in Computer Science, Vol. 10201)*. 696–723. doi:[10.1007/978-3-662-54434-1_26](https://doi.org/10.1007/978-3-662-54434-1_26)
- [35] Robert Krebbers, Amin Timany, and Lars Birkedal. 2017. Interactive proofs in higher-order concurrent separation logic. In *Principles of Programming Languages (POPL)*. doi:[10.1145/3009837.3009855](https://doi.org/10.1145/3009837.3009855)
- [36] Robbert Krebbers, Luko van der Maas, and Enrico Tassi. 2025. Inductive Predicates via Least Fixpoints in Higher-Order Separation Logic. In *Interactive Theorem Proving (ITP)*. doi:[10.4230/LIPICS.ITP.2025.27](https://doi.org/10.4230/LIPICS.ITP.2025.27)
- [37] Daan Leijen. 2014. Koka: Programming with Row Polymorphic Effect Types. In *Workshop on Mathematically Structured Functional Programming (MSFP)*, Vol. 153. 100–126. doi:[10.4204/EPTCS.153.8](https://doi.org/10.4204/EPTCS.153.8)
- [38] Xavier Leroy, Damien Doligez, Alain Frisch, Jacques Garrigue, Didier Rémy, KC Sivaramakrishnan, and Jérôme Vouillon. 2025. The OCaml system: Documentation and user’s manual. <https://ocaml.org/manual/5.3/index.html>
- [39] Sam Lindley, Conor McBride, and Craig McLaughlin. 2017. Do Be Do Be Do. In *Principles of Programming Languages (POPL)*. doi:[10.1145/3009837.3009897](https://doi.org/10.1145/3009837.3009897)
- [40] Craig McLaughlin. 2020. *Relational reasoning for effects and handlers*. Ph.D. Dissertation. University of Edinburgh, UK. doi:[10.7488/ERA/537](https://doi.org/10.7488/ERA/537)

- [41] Hiroshi Nakano. 2000. A Modality for Recursion. In *Logic in Computer Science (LICS)*. 255–266. doi:[10.1109/LICS.2000.855774](https://doi.org/10.1109/LICS.2000.855774)
- [42] Peter W. O’Hearn, John C. Reynolds, and Hongseok Yang. 2001. Local Reasoning about Programs that Alter Data Structures. In *Computer Science Logic (Lecture Notes in Computer Science, Vol. 2142)*. 1–19. doi:[10.1007/3-540-44802-0_1](https://doi.org/10.1007/3-540-44802-0_1)
- [43] Andrew Pitts and Ian Stark. 1999. *Operational reasoning for functions with local state*. Cambridge University Press, 227–274.
- [44] Gordon D. Plotkin and Matija Pretnar. 2008. A Logic for Algebraic Effects. In *Logic in Computer Science (LICS)*. 118–129. doi:[10.1109/LICS.2008.45](https://doi.org/10.1109/LICS.2008.45)
- [45] Gordon D. Plotkin and Matija Pretnar. 2009. Handlers of Algebraic Effects. In *European Symposium on Programming (ESOP) (Lecture Notes in Computer Science, Vol. 5502)*. 80–94. doi:[10.1007/978-3-642-00590-9_7](https://doi.org/10.1007/978-3-642-00590-9_7)
- [46] Gordon D. Plotkin and Matija Pretnar. 2013. Handling Algebraic Effects. *Logical Methods in Computer Science* 9, 4 (2013). doi:[10.2168/LMCS-9\(4:23\)2013](https://doi.org/10.2168/LMCS-9(4:23)2013)
- [47] Matija Pretnar. 2015. An Introduction to Algebraic Effects and Handlers. In *Mathematical Foundations of Programming Semantics (Electronic Notes in Theoretical Computer Science, Vol. 319)*. 19–35. doi:[10.1016/JENTCS.2015.12.003](https://doi.org/10.1016/JENTCS.2015.12.003)
- [48] John C. Reynolds. 2002. Separation Logic: A Logic for Shared Mutable Data Structures. In *Logic in Computer Science (LICS)*. 55–74. doi:[10.1109/LICS.2002.1029817](https://doi.org/10.1109/LICS.2002.1029817)
- [49] Alex Simpson and Niels Voorneveld. 2019. Behavioural Equivalence via Modalities for Algebraic Effects. *ACM Transactions on Programming Languages and Systems* 42 (2019). doi:[10.1145/3363518](https://doi.org/10.1145/3363518)
- [50] Simon Spies, Lennard Gähler, Joseph Tassarotti, Ralf Jung, Robbert Krebbers, Lars Birkedal, and Derek Dreyer. 2022. Later credits: resourceful reasoning for the later modality. *Proceedings of the ACM on Programming Languages* 6, ICFP (2022), 283–311. doi:[10.1145/3547631](https://doi.org/10.1145/3547631)
- [51] Wenhao Tang, Daniel Hillerström, Sam Lindley, and Garrett J. Morris. 2024. Soundly Handling Linearity. *Proceedings of the ACM on Programming Languages* 8, POPL (2024), 1600–1628. doi:[10.1145/3632896](https://doi.org/10.1145/3632896)
- [52] Wenhao Tang, Leo White, Stephen Dolan, Daniel Hillerström, Sam Lindley, and Anton Lorenzen. 2025. Modal Effect Types. *Proceedings of the ACM on Programming Languages* 9, OOPSLA (2025). doi:[10.1145/3720476](https://doi.org/10.1145/3720476)
- [53] The Iris Team. 2025. Iris Fixpoint Operators. https://gitlab.mpi-sws.org/iris/iris/-/blob/master/iris/bi/lib/fixpoint_mono.v
- [54] The Rocq Prover development team. 2025. *The Rocq Prover*. <https://rocq-prover.org/>
- [55] Amin Timany and Lars Birkedal. 2019. Mechanized Relational Verification of Concurrent Programs with Continuations. *Proceedings of the ACM on Programming Languages* 3, ICFP (2019), 105:1–105:28. doi:[10.1145/3341709](https://doi.org/10.1145/3341709)
- [56] Aaron Turon, Derek Dreyer, and Lars Birkedal. 2013. Unifying refinement and Hoare-style reasoning in a logic for higher-order concurrency. In *International Conference on Functional Programming (ICFP)*. 377–390. doi:[10.1145/2500365.2500600](https://doi.org/10.1145/2500365.2500600)
- [57] Orpheas van Rooij and Robbert Krebbers. 2025. Affect: An Affine Type and Effect System. *Proceedings of the ACM on Programming Languages* 9, POPL, 126–154. doi:[10.1145/3704841](https://doi.org/10.1145/3704841)
- [58] Simon Friis Vindum and Lars Birkedal. 2021. Contextual refinement of the Michael-Scott queue. In *Certified Programs and Proofs (CPP)*. 76–90. doi:[10.1145/3437992.3439930](https://doi.org/10.1145/3437992.3439930)
- [59] Simon Friis Vindum, Dan Frumin, and Lars Birkedal. 2022. Mechanized verification of a fine-grained concurrent queue from Meta’s Folly library. In *Certified Programs and Proofs (CPP)*. 100–115. doi:[10.1145/3497775.3503689](https://doi.org/10.1145/3497775.3503689)
- [60] Yizhou Zhang and Andrew C. Myers. 2019. Abstraction-safe effect handlers via tunneling. *Proceedings of the ACM on Programming Languages* 3, POPL (2019), 5:1–5:29. doi:[10.1145/3290318](https://doi.org/10.1145/3290318)
- [61] Yizhou Zhang, Guido Salvaneschi, Quinn Beightol, Barbara Liskov, and Andrew C. Myers. 2016. Accepting blame for safe tunneled exceptions. In *Programming Language Design and Implementation (PLDI)*. 281–295. doi:[10.1145/2908080.2908086](https://doi.org/10.1145/2908080.2908086)

Received 2025-07-10; accepted 2025-11-06